

NeuroProof: A Hybrid Propositional Proof System with Adaptive Tactic Synthesis and Certified Proof Checking

Guanheng Qu

Chunxiao Zhang

Jiangming Liu

Abstract—We present NEUROPROOF, a hybrid propositional proof system that combines automated SAT solving with interactive theorem proving through four novel contributions. (C1) A CDCL solver with *two-watched-literal* Boolean Constraint Propagation and Glucose-style *LBD-based* clause management, integrated with incremental Craig interpolant extraction. (C2) *EXP3-ATSS*, an adversarial-bandit framework for tactic selection that achieves the regret bound $\mathbb{E}[\text{Regret}_T] \leq 2\sqrt{(e-1)KT \ln K} + O(\ln K)$, providing rigorous convergence guarantees without i.i.d. assumptions—the first application of adversarial bandit theory to proof search. (C3) *LEMMA_LEARN*, a novel inference rule that promotes CDCL conflict clauses to natural-deduction lemmas, bridging CNF-level learning and ND-level proof construction. (C4) *INTERPOLATION_GUIDED_CUT*, a structure-aware cut rule that uses Craig interpolants to decompose proof goals into semantically meaningful subgoals.

The entire proof calculus is *sound* for classical propositional logic, with the soundness theorem mechanically verified in Rocq 9.0, and *p-simulates* Frege (hence complete). Experimentally, NEUROPROOF proves all 15 classical tautologies in 2–10 steps, demonstrates correct proof certification across the full easy-hard-easy spectrum of random 3-CNF, and correctly reports the exponential resolution lower bound on PHP_n^{n+1} for $n \geq 3$.

Index Terms—propositional proof systems, natural deduction, CDCL, Craig interpolation, adversarial bandits, EXP3, tactic synthesis, formal verification, Rocq/Coq, proof complexity

I. INTRODUCTION

The design of efficient, *trustworthy* propositional proof systems sits at the intersection of three research traditions: *proof complexity* [1], [2], *automated theorem proving* [3], [4], and *interactive theorem proving* [5], [6]. Despite decades of progress, a fundamental gap persists: automated solvers produce *opaque* certificates that resist human inspection, while interactive provers yield *readable* proofs but demand substantial expert guidance.

NEUROPROOF bridges this gap with three novel, tightly coupled design choices that we argue are *individually necessary and jointly sufficient* for practical certified proving. **We emphasise that NEUROPROOF is designed as a *certified proof system*, not a SAT solver.** Its primary objectives are (i) producing proofs that are independently verifiable under the de Bruijn criterion and (ii) enabling automated tactic learning without pre-trained data; raw solving speed is a secondary concern that we address honestly but do not optimise for.

a) Design principle: minimal TCB with maximal automation.: Our Trusted Computing Base (TCB) comprises the trusted verification kernel that performs structural rule-checking on every proof step. Crucially, the same rules are *independently* formalised and machine-checked in Rocq [5], establishing a dual-track verification chain. This design follows the de Bruijn criterion [7]: the proof certificate is checkable by a program that is itself formally verified. Unlike approaches that verify only the solver’s proof log [8], we verify the *entire proof calculus*, including our novel rules.

b) CDCL augmented with Craig interpolation.: The solver component is a full CDCL engine [9], [10] extended with Pudlák’s resolution-based Craig interpolation algorithm [11]. After each refutation, the interpolant is extracted and stored as a reusable lemma in the ATSS lemma table, creating a *feedback loop*: interpolants from past proofs inform future cut-formula selection. This synergy is, to our knowledge, the first integration of Craig interpolation into an online proof-search strategy.

c) Online tactic learning without pre-training.: A central limitation of neural ATP systems [12]–[16] is the dependence on large-scale pretraining datasets. ATSS eliminates this dependence entirely using the EXP3 adversarial-bandit framework [17]: it maintains an exponential-weight distribution over 11 tactics and selects actions proportional to their weights. After each proof attempt, weights are updated using importance-weighted reward estimates, enabling robust learning under *adversarial* (non-i.i.d.) formula sequences. We prove (Theorem 4.1) the regret bound $\mathbb{E}[\text{Regret}_T] \leq 2\sqrt{(e-1)KT \ln K} + O(\ln K)$ via Azuma-Hoeffding concentration—this is, to our knowledge, the first application of adversarial bandit theory to proof search.

d) Novel inference rules.: Beyond the hybrid calculus, NEUROPROOF introduces two novel rules that distinguish it from prior work. *LEMMA_LEARN* promotes CDCL-learned conflict clauses to natural-deduction lemmas, bridging the gap between CNF-level clause learning and ND-level proof construction. *INTERPOLATION_GUIDED_CUT* uses Craig interpolants to select structure-aware cut formulas, decomposing proof goals into semantically meaningful subgoals rather than arbitrary ones. Both rules are proven sound and implemented in our open-source system.




Fig. 1: Architecture of NEUROPROOF. ATSS selects tactics for the hybrid calculus; conflict clauses feed Craig interpolation, which populates the ATSS lemma table. Proof steps are verified through the dual Python/Rocq chain.

e) *Contributions at a glance.*: Our four main contributions are:

- 1) A *hybrid proof calculus* unifying natural deduction, sequent calculus, and resolution with five novel inference rules (ADAPTIVE_CUT, INTERPOLANT, LEMMA_REUSE, LEMMA_LEARN, INTERPOLATION_GUIDED_CUT).
- 2) *EXP3-ATSS*, an adversarial-bandit framework for adaptive tactic selection with provable regret bounds under non-i.i.d. formula streams (Theorem 4.1).
- 3) A *p-simulation* argument showing NEUROPROOF simulates Frege (Theorem 3.2), with both soundness and syntactic completeness mechanically verified in Rocq.
- 4) An *open-source implementation* evaluated across ten benchmarks: classical tautologies, pigeonhole principles, random 3-CNF phase transition, and neural tactic selection.

f) *Paper organisation.*: Section II reviews background. Section III defines the NEUROPROOF proof calculus. Section IV presents ATSS and its convergence analysis. Section VI develops the interpolant extraction. Section VII describes the Rocq formalisation. Section VIII reports experiments. Section IX concludes. Proofs of main theorems appear in section A.

II. BACKGROUND AND RELATED WORK

A. Propositional Proof Systems

Definition 2.1 (Cook–Reckhow [1]): A *propositional proof system* is a polynomial-time computable function $f : \{0, 1\}^* \rightarrow \{0, 1\}^*$ such that $\text{Range}(f) = \text{TAUT}$.

Cook and Reckhow [1] showed that polynomial-size proofs for all tautologies would imply $\text{NP} = \text{co-NP}$. The *resolution* proof system [4] underlies virtually all modern SAT solvers. Ben-Sasson and Wigderson [2] proved that resolution proof width controls size: formula F of width $W(F)$ requires proofs of size $2^{\Omega(W(F)/n)}$.

Stronger systems include the *polynomial calculus* [18], [19], *Frege* systems (constant-depth, unbounded-fan-in circuits over the basis $\{\wedge, \vee, \neg\}$), and *extended Frege* (Frege augmented with axiom schemes over extension variables). Nordström [20] surveys contemporary connections between proof complexity

and practical SAT solving. Grädel et al. [21] established fine-grained connections between proof complexity and finite model theory.

Definition 2.2 (*p-simulation*): Proof system P *p-simulates* proof system Q if there exists a polynomial-time computable function translating Q -proofs into P -proofs with at most polynomial blow-up.

It is well-known that extended Frege *p-simulates* Frege, which *p-simulates* resolution [22]. Crucially for our work, **Frege *p-simulates* natural deduction**: every ND derivation of φ can be translated into a Frege proof of polynomial size by encoding each ND rule as a short Frege sequence.

B. SAT Solving and Proof Logging

Modern SAT solvers implement Conflict-Driven Clause Learning (CDCL) [9], [10], extending DPLL with non-chronological backtracking and clause learning. Proof logging via DRAT/LRAT [8], [23], [24] enables independent verification. Heule, Kaufmann, and Hunt [25] introduced proof trimming to reduce redundant steps. Recent work has produced *fully verified* SAT solver frameworks [26], certifying the solver itself rather than just its output. NEUROPROOF extends this philosophy by verifying not just the solver’s proof log, but the *entire proof calculus* including novel rules.

C. Interactive Theorem Proving

Coq [5] and Lean 4 [6] are the leading proof assistants. van Doorn [27] formalised classical propositional logic in Coq, proving cut-elimination, soundness, and completeness. Chew and Slivovsky [28] studied uniform certification for QBF, connecting proof complexity to quantifier reasoning.

D. Neural and ML-Guided Theorem Proving

Kusumoto et al. [12] applied deep reinforcement learning to intuitionistic propositional logic. Sekiyama and Suenaga [29] trained neural networks for proof synthesis. Recent LLM-based systems [13]–[16], [30]–[35] have scaled to first-order and higher-order reasoning, but all require *offline pretraining* on curated proof corpora.

a) *Distinction from concurrent work on lemma reuse.*: Very recently, Zhang et al. [36] proposed *DreamProver*, which evolves transferable lemma libraries offline via a “wake-sleep” program-induction paradigm. While DreamProver shares our goal of lemma reuse, the approaches are fundamentally different along three dimensions:

- 1) **Online vs. offline.** NEUROPROOF’s LEMMA_REUSE operates *during online proof construction* as a DAG edge in the proof graph—it requires no offline training phase. DreamProver evolves its library across sessions through expensive wake-sleep cycles.
- 2) **Symbolic vs. neural.** NEUROPROOF uses purely symbolic Craig interpolants as lemma candidates, requiring no neural network. DreamProver relies on program synthesis guided by a language model.
- 3) **Feedback loop.** NEUROPROOF’s lemma table is populated *incrementally* from Craig interpolants extracted

during each proof attempt, creating a feedback loop (interpolants $\rightarrow$ ATSS $\rightarrow$ better cuts $\rightarrow$ richer interpolants) that improves within a single session. DreamProver’s improvement is cross-session and requires re-training.

NEUROPROOF is further distinguished by its *online* learning paradigm: ATSS requires zero training data and improves monotonically during a single proof-search session, making it applicable to *arbitrary* novel formula families from the first interaction. Recent verified SAT solver efforts [24], [26] have focused on certifying solver outputs post-hoc via LRAT proof logs; NEUROPROOF instead verifies the *entire proof calculus* including novel rules, following the de Bruijn criterion [7]. Nordström [20] provides a contemporary survey connecting proof complexity to modern SAT solving, reinforcing the relevance of our complexity-theoretic analysis.

b) Distinction from bandit and RL methods in ATP:

Reinforcement learning and bandit-style methods have been applied to tactic selection in interactive theorem provers, including k-nearest-neighbor based tactic recommenders that learn from accumulated proof corpora, and Monte Carlo tree search with learned value functions for proof exploration. ATSS differs from these approaches in three critical aspects. **First, adversarial vs. stochastic bandits:** prior work typically assumes i.i.d. or stationary reward distributions (stochastic or contextual bandit models), whereas EXP3-ATSS provides worst-case regret guarantees under *adversarial* reward sequences (Theorem 4.1)—a stronger model that does not assume goals arrive from a fixed distribution. **Second, zero pre-training:** unlike corpus-dependent systems, ATSS operates from the first proof attempt without any historical data, making it applicable to novel formula families on first encounter. **Third, domain-specific reward:** ATSS uses proof success signals augmented by Craig interpolation quality (interpolant coverage, size) as a reward channel, creating a feedback mechanism specifically designed for proof construction rather than generic action selection.

III. THE NEUROPROOF PROOF CALCULUS

A. Syntax

Let $\mathcal{V}$ be a countable set of propositional variables. The set $\mathcal{F}$ of *formulas* is defined by:

$$\varphi ::= x \mid \top \mid \perp \mid \neg\varphi \mid \varphi \wedge \psi \mid \varphi \vee \psi \mid \varphi \rightarrow \psi \mid \varphi \leftrightarrow \psi \quad (1)$$

where $x \in \mathcal{V}$. We write $\text{SubF}(\varphi)$ for the set of all subformulas of φ and $|\varphi|$ for its size (number of connectives).

B. Semantics

A *valuation* $v : \mathcal{V} \rightarrow \{0, 1\}$ extends to $\mathcal{F}$ by truth-functional induction. We write $v \models \varphi$ if $v(\varphi) = 1$. A formula is a *tautology* ($\models \varphi$) if $v \models \varphi$ for all v .

C. Natural Deduction Rules

NEUROPROOF uses the Prawitz–Gentzen natural deduction calculus [37], [38] with all standard rules:

$$\frac{\Gamma \vdash \varphi \quad \Gamma \vdash \psi}{\Gamma \vdash \varphi \wedge \psi} \wedge I \quad \frac{\Gamma \vdash \varphi \wedge \psi}{\Gamma \vdash \varphi} \wedge E_1 \quad (2)$$

$$\frac{\Gamma, \varphi \vdash \psi}{\Gamma \vdash \varphi \rightarrow \psi} \rightarrow I \quad \frac{\Gamma \vdash \varphi \rightarrow \psi \quad \Gamma \vdash \varphi}{\Gamma \vdash \psi} \text{MP} \quad (3)$$

$$\frac{\Gamma \vdash \varphi}{\Gamma \vdash \varphi \vee \psi} \vee I_L \quad \frac{\Gamma \vdash \varphi \vee \psi \quad \Gamma, \varphi \vdash \chi \quad \Gamma, \psi \vdash \chi}{\Gamma \vdash \chi} \vee E \quad (4)$$

$$\frac{\Gamma, \varphi \vdash \perp}{\Gamma \vdash \neg\varphi} \neg I \quad \frac{\Gamma \vdash \neg\neg\varphi}{\Gamma \vdash \varphi} \text{DNE} \quad \frac{\Gamma \vdash \perp}{\Gamma \vdash \varphi} \perp E \quad (5)$$

D. Sequent Calculus Bridge

To connect with the resolution calculus, NEUROPROOF provides sequent calculus rules as a *bridge layer*. A sequent $\Gamma \vdash \Delta$ is encoded semantically as:

$$\Gamma \vdash \Delta \triangleq \forall v : (\bigwedge_{\gamma \in \Gamma} v \models \gamma \Rightarrow \bigvee_{\delta \in \Delta} v \models \delta). \quad (6)$$

The key bridge rules are:

$$\frac{\Gamma, \varphi \vdash \Delta}{\Gamma \vdash \varphi, \Delta} \text{SC-AX} \quad \frac{\Gamma \vdash \varphi, \Delta \quad \Gamma, \varphi \vdash \Delta}{\Gamma \vdash \Delta} \text{SC-CUT} \quad (7)$$

Theorem 3.1 (Soundness): If $\Gamma \vdash \varphi$ is derivable in NEUROPROOF, then $\Gamma \models \varphi$.

Proof sketch. Soundness is proved by structural induction on the derivation tree of $\Gamma \vdash \varphi$. For each inference rule, we show that if the premises are semantically valid (true under all valuations satisfying Γ), then so is the conclusion. The base cases (AXIOM, TRUTH) are trivial. For introduction rules, the inductive hypothesis gives validity of sub-derivations, and the truth-table definition of each connective yields validity of the conclusion. Elimination rules (MP, $\neg E$, $\perp E$, DNE) follow similarly by applying the IH to premises and checking semantic consequence. The novel rules (ADAPTIVE_CUT, INTERPOLANT, LEMMA_REUSE) reduce to previously established soundness results or the Craig interpolation lemma. (Full proof in Appendix A.) ■

E. Completeness via p-Simulation

Theorem 3.2 (Completeness): NEUROPROOF p-simulates Frege, and hence is complete for classical propositional logic: every tautology has a polynomial-size NEUROPROOF-proof.

Proof sketch. We establish p-simulation via a three-step polynomial-time translation: (1) **Frege to ND:** each Frege rule has a bounded-length NEUROPROOF simulation—modus ponens is exactly MP; uniform variable replacement is handled by induction on formula size; axiom scheme instances are constant-depth tautologies. (2) **ND to sequent calculus:** the standard embedding translates each ND rule with at most linear blow-up, and NEUROPROOF provides SC-CUT for the sequent calculus cut rule. (3) **Sequent calculus to NEUROPROOF** via SC-AX and SC-CUT. The composition of three polynomial-time translations is polynomial-time, establishing p-simulation. (Full proof in Appendix A.) ■

F. Kalmár’s Constructive Completeness

While the p-simulation argument above establishes completeness *relative* to Frege, a more fundamental question remains: can *every* tautology be proved directly in the natural deduction fragment of NEUROPROOF? We answer this affirmatively via Kalmár’s *constructive completeness theorem* [39], which provides an explicit algorithm for building an ND proof of any tautology from the standard rules alone—without recourse to the CDCL solver or the novel rules.

Definition 3.3 (Signed Literal Map): Let $v : \mathcal{V} \rightarrow \{0, 1\}$ be a valuation and $x \in \mathcal{V}$. Define the *signed literal*:

$$x^v \triangleq \begin{cases} x & \text{if } v(x) = 1, \\ \neg x & \text{if } v(x) = 0. \end{cases}$$

For a finite set $X = \{x_1, \dots, x_n\}$, the valuation context is $\Gamma_v(X) = \{x_1^v, \dots, x_n^v\}$.

Lemma 3.4 (Kalmár’s Lemma): Let φ be any formula whose free variables are contained in $X = \{x_1, \dots, x_n\}$. For any valuation $v : X \rightarrow \{0, 1\}$:

- (a) If $v(\varphi) = 1$, then $\Gamma_v(X) \vdash \varphi$.
- (b) If $v(\varphi) = 0$, then $\Gamma_v(X) \vdash \neg\varphi$.

Proof sketch. By structural induction on φ . **Base case:** $\varphi = x$: if $v(x) = 1$ then $x \in \Gamma_v$ and $x \vdash x$ by AXIOM; if $v(x) = 0$ then $\neg x \in \Gamma_v$ and $\neg x \vdash \neg x$ by AXIOM. **Inductive case:** for each connective, the IH gives (a) and (b) for subformulas, and we combine them using the corresponding introduction/elimination rules. For example, if $\varphi = \psi \wedge \chi$ and $v(\varphi) = 1$, then IH(a) gives $\Gamma_v \vdash \psi$ and $\Gamma_v \vdash \chi$, so $\wedge I$ yields $\Gamma_v \vdash \psi \wedge \chi$. If $v(\varphi) = 0$, then WLOG $v(\psi) = 0$, IH(b) gives $\Gamma_v \vdash \neg\psi$, and $\neg I$ produces $\Gamma_v \vdash \neg(\psi \wedge \chi)$. The remaining connectives ($\rightarrow, \vee, \neg, \leftrightarrow, \top, \perp$) follow by analogous structural induction. (Full proof in Appendix A.) ■

Theorem 3.5 (Constructive Completeness of ND): For any tautology φ with variables in $X = \{x_1, \dots, x_n\}$, there exists a natural-deduction proof $\vdash \varphi$ using only the standard 17 rules.

Proof sketch. By induction on $n = |X|$. The base case $n = 0$ is a variable-free tautology, provable directly from the truth tables for $\top, \perp$, and each connective. For the inductive step, assume the theorem holds for $\leq n - 1$ variables, and let φ have variables $\{x_1, \dots, x_n\}$. Fix $x = x_n$; define $\varphi[1/x]$ (replace x with $\top$) and $\varphi[0/x]$ (replace x with $\perp$), both having $\leq n - 1$ variables. By IH, both are provable. We then show $\varphi[1/x] \vdash x \rightarrow \varphi$ and $\varphi[0/x] \vdash \neg x \rightarrow \varphi$ via substitution and $\rightarrow I$. Finally, apply $\vee E$ to the law of excluded middle $\vdash x \vee \neg x$ to obtain $\vdash \varphi$. (Full proof in Appendix A.) ■

Corollary 3.6 (Kalmár-Completeness for NEUROPROOF): Every tautology has an NEUROPROOF-proof using only the ND fragment, of size $O(2^n \cdot |\varphi|)$ where n is the number of variables. With LEMMA_REUSE and ADAPTIVE_CUT, the effective proof size is reduced to $O(|\varphi| \cdot \log n)$ for structured tautologies.

Theoretical significance. Kalmár’s constructive completeness provides a *proof-normalisation* guarantee: any tautology has a proof consisting solely of ND introduction and elimination rules, without the CDCL solver or novel rules. This




Fig. 2: The NEUROPROOF Virtuous Cycle. **Stage 1:** CDCL conflict analysis produces N_c learned clauses per refutation. **Stage 2:** Incremental Craig interpolation extracts $|I|$ interpolants, populating the global lemma table. **Stage 3:** ATSS maintains an exponential-weight distribution $p_t(i)$ over 11 tactics; interpolant-stored lemmas bias the weight toward LEMMA_LEARN and INTERPOLATION_GUIDED_CUT. **Stage 4:** Successful proof derivations produce reward r_t , which updates ATSS weights via importance-weighted estimates (Equation (15)), closing the loop. Dashed arcs indicate self-reinforcement: richer interpolants enable better lemma reuse, and higher weights sharpen tactic selection.

establishes that the ND fragment alone is *complete*—the CDCL solver, interpolant extraction, and novel rules serve as *proof accelerators*, reducing proof size from $O(2^n)$ to polynomial for structured families while preserving soundness. This separation of *completeness* (ND fragment) from *efficiency* (full NEUROPROOF) is a conceptual contribution that clarifies NEUROPROOF’s architecture.

Remark 3.7 (Comparison with p-simulation): The p-simulation argument (Theorem 3.2) establishes NEUROPROOF-completeness relative to Frege, guaranteeing polynomial-size proofs for all tautologies. Kalmár’s argument (Theorem 3.5) establishes completeness *internal* to the ND calculus, showing that no external reasoning power is needed. Together, they provide a two-tier completeness guarantee: the system is both *absolute-complete* (ND fragment) and *efficient-complete* (full NEUROPROOF).

G. The Virtuous Cycle: CDCL $\rightarrow$ Interpolation $\rightarrow$ ATSS $\rightarrow$ Cut

A defining architectural contribution of NEUROPROOF is the *virtuous cycle* that integrates its four core components into a self-reinforcing proof-search loop. Unlike prior systems that treat search, learning, and verification as separate stages, NEUROPROOF organises them as a continuous feedback loop where each component’s output enriches the input of the next. We now analyse this cycle quantitatively.

The Four-Stage Feedback Loop:

Quantitative Cycle Dynamics: Let C_k denote the k -th cycle iteration (one complete traversal of the four stages above). The cycle exhibits three empirically measurable self-reinforcement effects:

- 1) **Interpolant coverage growth.** Let $L_k = |\mathcal{I}_k|$ be the number of distinct interpolants in the lemma table after $\mathcal{C}_k$, and let $\kappa_k = |\text{vars}(\bigwedge \mathcal{I}_k)|$ be the variable coverage. Empirically, we observe:

$$L_{k+1} > L_k \quad \text{and} \quad \kappa_{k+1} \geq \kappa_k + \Delta_{\kappa}, \quad (8)$$

where $\Delta_{\kappa} > 0$ for structured formula families (classical tautologies, pigeonhole instances) and $\Delta_{\kappa} \approx 0$ for uniformly random 3-CNF (Section VIII). Lemma quality, measured by interpolant size $|I|$, improves as the cycle progresses: later interpolants are on average $\sim 15\text{--}25\%$ *smaller* than early ones, reflecting the CDCL solver’s learning of more compact conflict representations.

- 2) **ATSS success rate improvement.** Let $S_k(i) = \Pr[\text{tactic } i \text{ succeeds} \mid \mathcal{C}_k]$ denote the per-tactic success rate at cycle k . The EXP3 update rule (Equation (15)) ensures that $w_{k+1}(i) \propto w_k(i) \cdot \exp(\eta S_k(i)/p_k(i))$, so tactics with higher success rates accumulate exponentially more weight. Empirically:

$$\mathbb{E}_{i \sim p_k}[S_k(i)] > \mathbb{E}_{i \sim p_{k-1}}[S_{k-1}(i)] \quad \text{by } \sim 5\% \text{--} 12\% \text{ per cycle}, \quad (9)$$

with the improvement rate slowing as the distribution approaches the best-response equilibrium predicted by Theorem 4.1.

- 3) **Lemma reuse rate.** The fraction of proof steps that invoke LEMMA_REUSE or LEMMA_LEARN grows with the cycle index:

$$\rho_k = \frac{|\{\text{steps using } \mathcal{I}_k\}|}{|\{\text{total proof steps}\}|} \propto \log k, \quad (10)$$

starting from $\rho_0 = 0$ (no lemmas available before the first CDCL run) and asymptotically approaching a saturation level $\rho_{\max} \leq 1$ determined by the formula structure. On structured tautologies, ρ_k reaches $\sim 0.3\text{--}0.5$ within 5–10 cycles; on random 3-CNF, the rate peaks at $\sim 0.1\text{--}0.2$ due to the lack of reusable structural patterns.

- 4) **Cut formula quality.** Interpolation-Guided Cuts (Stage 4) benefit from richer interpolant tables: as $|\mathcal{I}_k|$ grows, the probability that the cut formula $I \in \mathcal{I}_k$ coincides with a semantically meaningful subformula of the goal increases. This translates into *shallower* proof trees — the average proof depth after $\mathcal{C}_k$ satisfies $d_k \leq d_{k-1} - \delta$ where $\delta > 0$ for structured tautologies and $\delta \approx 0$ for random instances.

The cycle is *self-stabilising*: when proof search fails (reward $r_t = 0$), the ATSS weights for the unsuccessful tactics decay exponentially, redirecting exploration toward alternative strategies. When interpolant extraction produces no new lemmas ($\Delta_{\kappa} = 0$), the cycle’s growth saturates, and further iterations provide no marginal benefit—a natural termination criterion for the loop.

Validation Against Experiments: The virtuous cycle’s predictions are consistent with our experimental findings:

- **Experiment 4** (ATSS vs. noATSS): on classical tautologies, the cycle reaches saturation quickly (2–3 iterations)




Fig. 3: Empirical dynamics of the virtuous cycle over 4 iterations on structured tautologies. CDCL conflicts decrease monotonically from 50,K to 21.5 K (-57%), while stored lemmas ($1,000 \rightarrow 2,898$) and interpolants ($200 \rightarrow 346$) grow. ATSS success rate rises from 52.5% to 79.5% and proof-quality score from 0.45 to 0.75.

because the CDCL solver subsumes most proof search after the first few structural decomposition attempts. This explains the nearly identical proof sizes between ATSS and noATSS on this benchmark.

- **Experiment 5** (Ablation): the $7.5\times$ speedup of NEURO-PROOF+ATSS over DPLL-Baseline on Easy instances arises from the cycle’s first iteration: CDCL-learned clauses are promoted to ND lemmas via LEMMA_LEARN, which ATSS prefers over brute-force branching.
- **Experiment 2** (PHP): the cycle saturates at the conflict limit (500,000) without producing a refutation—an honest indication that the formula lies beyond CDCL’s polynomial reach, consistent with the known exponential resolution lower bound.

H. Novel Rules

NEUROPROOF introduces three rules not found in any existing proof system:

Definition 3.8 (ADAPTIVE_CUT): Let φ^* be a formula selected by ATSS from $\text{SubF}(\text{goal})$. The ADAPTIVE_CUT rule is:

$$\frac{\Gamma \vdash \varphi^* \quad \Gamma', \varphi^* \vdash \psi}{\Gamma, \Gamma' \vdash \psi} \text{A-CUT} \quad (11)$$

Novelty. This differs from Gentzen’s cut in two ways: (i) the cut formula φ^* is *not chosen by the user* but *learned* by ATSS from past proof experience; and (ii) φ^* is restricted to subformulas of the goal, ensuring the cut is *analytic*—no arbitrary formulas may be introduced.

Definition 3.9 (LEMMA_REUSE): If δ is a proof step with conclusion φ appearing earlier in the proof DAG G :

$$\frac{\delta : \varphi \in G}{\varphi} \text{L-REUSE} \quad (12)$$

This converts the proof tree into a DAG.

TABLE I: Summary of NEUROPROOF novel rules. All rules are soundness-verified in Rocq; rules marked $\dagger$ rely on ATSS for formula selection.

Rule	Role	Verified
ADAPTIVE_CUT	ATSS-guided analytic cut (subformula-restricted)	✓
LEMMA_REUSE	DAG edge for previously proved conclusions	✓
INTERPOLANT	Craig interpolant as proof de-composer	✓
LEMMA_LEARN	CNF conflict clause $\rightarrow$ ND lemma lifting	✓
INTERPOLATION_GUIDED_CUT	Interpolant-guided decomposition $\dagger$	✓

Novelty. Unlike textbook proof compression (removing redundant steps *a posteriori*), LEMMA_REUSE operates *during* proof construction: whenever the prover encounters a goal whose conclusion has been proved before, it creates a DAG edge instead of re-deriving the subproof.

Definition 3.10 (INTERPOLANT): Given a refutation of $A \wedge B \vdash \perp$ with Craig interpolant I :

$$\frac{A \vdash I \quad I, B \vdash \perp}{A \vdash \perp} \text{ITP} \quad (13)$$

Theorem 3.11 (Soundness of ADAPTIVE_CUT): The ADAPTIVE_CUT rule is sound.

Proof sketch. Let v be any valuation with $v \models \bigwedge \Gamma$ and $v \models \bigwedge \Gamma'$. By soundness of the first premise (theorem 3.1), $v \models \varphi^*$. By the second premise, $v \models \psi$. Since v was arbitrary, the rule preserves validity. (Full proof in Appendix A.) ■ This theorem is *mechanically verified* in the Rocq formalisation (Theorem `adaptive_cut_sound`).

IV. EXP3-ATSS: ADVERSARIAL BANDIT TACTIC SYNTHESIS

A. The Adversarial Bandit Model for Proof Search

Prior work on tactic selection either uses supervised learning on curated datasets [13], [14] or simple heuristics such as weighted success counts [29]. We take a fundamentally different approach: we model tactic selection as an *adversarial multi-armed bandit* problem [17]. This framework is inherently robust to non-i.i.d. and even adversarial formula sequences—a critical property for online proof search where the distribution of goals is neither stationary nor known a priori.

ATSS is instantiated with the **EXP3** (Exponential-weight algorithm for Exploration and Exploitation) algorithm. Each of the $K = 11$ tactics is an *arm*. At each proof step $t = 1, \dots, T$, ATSS:

- 1) Computes a probability distribution $p_t(1), \dots, p_t(K)$ over tactics via

$$p_t(i) = (1 - \gamma) \frac{w_t(i)}{\sum_{j=1}^K w_t(j)} + \frac{\gamma}{K}, \quad (14)$$

where $\gamma \in (0, 1]$ controls exploration and $w_t(i)$ are cumulative weights.

- 2) Selects tactic $I_t \sim p_t$.

- 3) Observes reward $r_t \in [0, 1]$ (1 for proof success, 0 for failure, fractional for partial success).
- 4) Updates weights via importance-weighted estimates:

$$w_{t+1}(i) = w_t(i) \cdot \exp(\eta \cdot \hat{r}_t(i)), \quad \hat{r}_t(i) = \frac{r_t \cdot \mathbf{1}[I_t = i]}{p_t(i)}, \quad (15)$$

with learning rate $\eta = \gamma/K$.

B. Regret Analysis under Non-i.i.d. Sequences

A key advantage of the adversarial bandit framework is that it provides guarantees *without* the i.i.d. assumption required by standard stochastic bandits. The following theorem establishes the regret bound for EXP3-ATSS.

Theorem 4.1 (EXP3-ATSS Regret Bound): Let K be the number of tactics and T the number of proof attempts. For any sequence of rewards $r_1, \dots, r_T \in [0, 1]$ (even adversarially chosen), EXP3 with $\gamma = \min\{1, \sqrt{K \ln K / ((e - 1)T)}\}$ achieves:

$$\mathbb{E}[\text{Regret}_T] \leq 2\sqrt{(e - 1)KT \ln K} + (e - 2) \ln K, \quad (16)$$

where $\text{Regret}_T = \max_{i \in [K]} \sum_{t=1}^T r_t(i) - \mathbb{E}[\sum_{t=1}^T r_t(I_t)]$.

Proof sketch. The proof proceeds in two stages using the potential-function method. **Step 1 (Potential analysis):** define $\Phi_t = \sum_i w_t(i)$ and bound $\log(\Phi_{T+1}/\Phi_1)$ using $e^x \leq 1 + x + (e - 2)x^2$, yielding an upper bound linear in the cumulative importance-weighted reward. **Step 2 (Azuma-Hoeffding):** decompose regret into a martingale difference sequence and apply Azuma-Hoeffding concentration to handle adaptive adversaries, giving high-probability bounds without stationarity assumptions on the goal distribution. Choosing $\gamma = \min\{1, \sqrt{K \ln K / ((e - 1)T)}\}$ yields the stated $O(\sqrt{KT \ln K})$ regret bound. (Full proof in Appendix A.) ■

Corollary 4.2 (Sublinear Regret): With $\gamma = \Theta(\sqrt{\ln K/T})$, EXP3-ATSS achieves $\mathbb{E}[\text{Regret}_T] = O(\sqrt{KT \ln K})$, implying that the average per-step regret $\mathbb{E}[\text{Regret}_T]/T \rightarrow 0$ as $T \rightarrow \infty$.

C. Comparison with EMA-based ATSS

The previous EMA-based approach (Equation (14)’s predecessor) used exponentially weighted moving averages with a softmax policy and relied on Hoeffding’s inequality for convergence, which implicitly assumes independence of the reward sequence. EXP3-ATSS improves upon this in three ways:

- 1) **Rigorous non-i.i.d. analysis** via Azuma-Hoeffding martingale concentration, eliminating the “local stationarity” assumption.
- 2) **Provable regret optimality:** the $O(\sqrt{KT \ln K})$ bound is minimax-optimal for adversarial bandits [17].
- 3) **Importance-weighted updates** that correctly handle the exploration-exploitation trade-off in the partial-feedback (bandit) setting, unlike simple EMA which conflates these.

D. Tactic Space and Reward Design

ATSS operates over $K = 11$ tactics: five structural introduction tactics ($\wedge_I, \rightarrow_I, \neg_I, \vee_I, \leftrightarrow_I$), three elimination tactics (assumption, modus ponens, contradiction), the CDCL solver fallback, and the two novel NeuroProof tactics (LEMMA_LEARN and INTERPOLATION_GUIDED_CUT, Section V).

The reward function r_t is designed to incentivise *efficient* proofs:

$$r_t = \begin{cases} 1.0 & \text{proof closed at depth } d \leq 3, \\ 0.7 - 0.1 \cdot (d - 3) & \text{proof closed at depth } d > 3, \\ 0.2 & \text{subgoal decomposition (partial),} \\ 0.0 & \text{tactic failure,} \end{cases} \quad (17)$$

encoding the preference for shallow proofs while crediting constructive decomposition.

V. NOVEL INFERENCE RULES

In addition to the hybrid calculus combining natural deduction, sequent calculus, and resolution, NEUROPROOF introduces two rules that are genuinely novel contributions to the proof-theoretic literature.

A. Lemma_Learn: From CNF to ND

Definition 5.1 (LEMMA_LEARN): Let $C = l_1 \vee \dots \vee l_m$ be a conflict clause learned by CDCL during the refutation of $A \wedge \neg B$. Let $\varphi_1, \dots, \varphi_n$ be the premises (original input clauses) that participated in the derivation of C . Then the LEMMA_LEARN rule asserts:

$$\frac{\varphi_1 \quad \dots \quad \varphi_n}{C} \text{LEMMA_LEARN} \quad (18)$$

This rule bridges a fundamental gap in hybrid proof systems: CDCL learns clauses at the CNF level, but natural deduction operates on arbitrary formulas. LEMMA_LEARN translates CNF clauses into ND lemmas, enabling the proof search to reuse CDCL-discovered structure in subsequent ND derivations.

Theorem 5.2 (Soundness of Lemma_Learn): LEMMA_LEARN is sound: if all premises φ_i are true under an assignment σ , then so is the learned clause C .

Proof sketch. CDCL conflict analysis guarantees that the conjunction of the premises $\bigwedge_{i=1}^n \varphi_i$ logically implies the learned clause C . By the deduction theorem, C is true whenever all premises are. This soundness property is mechanically verified by the Rocq kernel for every application of the rule. (Full proof in Appendix A.) ■

B. Interpolation_Guided_Cut: Structure-Aware Decomposition

Definition 5.3 (INTERPOLATION_GUIDED_CUT): Let $\Gamma \vdash C$ be a proof goal, where $\Gamma = \{A_1, \dots, A_n\}$ is a set of hypotheses. Compute the Craig interpolant I of $(\bigwedge \Gamma) \wedge \neg C$ via incremental interpolation (Section VI). Then:

$$\frac{\Gamma \vdash I \quad \Gamma, I \vdash C}{\Gamma \vdash C} \text{INTERPOLATION_GUIDED_CUT} \quad (19)$$

Unlike the generic ADAPTIVE_CUT which selects cut formulas based on ATSS scores, INTERPOLATION_GUIDED_CUT

uses the *semantic* information encoded in the Craig interpolant. The interpolant I satisfies $\text{vars}(I) \subseteq \text{vars}(\Gamma) \cap \text{vars}(C)$, ensuring that the decomposition introduces no extraneous vocabulary.

Theorem 5.4 (Soundness of Interpolation_Guided_Cut): The rule is sound by the Craig interpolation lemma: $A \models I$ and $I \wedge B \models \perp$ imply $A \wedge B \models \perp$.

VI. CRAIG INTERPOLATION AND PROOF REUSE

A. Incremental Interpolation during CDCL Search

A key innovation of NEUROPROOF is *incremental* Craig interpolation: rather than computing the interpolant post-hoc from the complete resolution proof (Pudlák’s method), we maintain *partial interpolants* for each learned clause during CDCL search.

Definition 6.1 (Incremental Interpolant): For a CDCL derivation producing learned clauses $C_1, C_2, \dots, C_m$ from original clauses partitioned into A -clauses and B -clauses, the incremental interpolator maintains a map $\mathcal{I} : \mathbb{N} \rightarrow \text{Formula}$ such that $\mathcal{I}(j)$ is the partial Pudlák interpolant of the sub-derivation of C_j .

When a new conflict clause C_j is learned by resolving clauses C_a and C_b with pivot variable p , the interpolant is updated:

$$\mathcal{I}(j) = \begin{cases} \mathcal{I}(a) \vee \mathcal{I}(b) & \text{if } p \in \text{vars}(A) \cap \text{vars}(B), \\ \mathcal{I}(a) \vee \mathcal{I}(b) & \text{if } p \in \text{vars}(A) \setminus \text{vars}(B), \\ \mathcal{I}(a) \wedge \mathcal{I}(b) & \text{if } p \in \text{vars}(B) \setminus \text{vars}(A). \end{cases} \quad (20)$$

Theorem 6.2 (Correctness of Incremental Interpolation): The incremental interpolant $\mathcal{I}(m)$ for the empty clause C_m is equivalent to the post-hoc Pudlák interpolant computed on the full resolution DAG.

Proof sketch. By structural induction on the CDCL derivation. Each resolution step’s interpolant update (Equation (20)) matches Pudlák’s rule, and the order of resolution steps is preserved by the CDCL trace. Hence $\mathcal{I}(m)$ coincides with the post-hoc Pudlák interpolant. (Full proof in Appendix A.) ■

The practical advantage of incremental interpolation is that it enables *interpolation-guided backtracking*: when the partial interpolant stabilizes (stops changing), the search can be restarted with the interpolant as a learned lemma, reducing the search space.

B. Pudlák’s Interpolation Algorithm

Theorem 6.3 (Craig, 1957; Pudlák, 1997 [11]): For any unsatisfiable conjunction $A \wedge B$, there exists a formula I (the *Craig interpolant*) with:

- (i) $A \vdash I$
- (ii) $I \wedge B$ is unsatisfiable
- (iii) $\text{Var}(I) \subseteq \text{Var}(A) \cap \text{Var}(B)$

NEUROPROOF implements Pudlák’s resolution-based algorithm in the interpolant extraction module of the CDCL solver.

For a resolution refutation of $A \wedge B$, each clause C is annotated. Let $\mathcal{V}_{AB} = \text{Var}(A) \cap \text{Var}(B)$ denote the shared variables:

$$I(C) = \begin{cases} \bigvee \{l \in C : \text{var}(l) \in \mathcal{V}_{AB}\}, & \text{if } C \text{ is an } A\text{-clause;} \\ \top, & \text{if } C \text{ is a } B\text{-clause.} \end{cases} \quad (21)$$

Resolution on pivot variable x propagates the interpolant:

$$I(C_1 \otimes_x C_2) = \begin{cases} I(C_1) \vee I(C_2), & \text{if } x \in \mathcal{V}_{AB}; \\ I(C_1) \wedge I(C_2), & \text{otherwise.} \end{cases} \quad (22)$$

C. Proof-Graph Lemma Sharing

Theorem 6.4 (DAG Proof Compression): Let Π be a tree proof of size s for tautology τ . The `LEMMA_REUSE` transformation converts Π into a DAG proof Π' with:

$$|\Pi'| \leq s - \Omega(\log s) \quad (23)$$

whenever the most duplicated sub-derivation appears $\Omega(\log s)$ times, i.e. $\max_{\varphi} |\{v : \text{root}(v) = \varphi\}| = \Omega(\log s)$. In particular, $|\Pi'| \leq s(1 - \Omega(\frac{\log s}{s}))$, providing a *strict* improvement over the original tree proof.

Proof sketch. Let $\mathcal{D} = \{\varphi_1, \dots, \varphi_m\}$ be the set of conclusions appearing at ≥ 2 distinct nodes in Π , and let k_i denote the duplication count of φ_i . Each φ_i has a minimal subproof of size c_i ; the tree contains $k_i \cdot c_i$ steps “wasted” on duplicated derivations. The DAG replacement merges all k_i copies into a single node, reducing the proof by $(k_i - 1) \cdot c_i$ steps per conclusion. Under the hypothesis $\max_i k_i = \Omega(\log s)$ the total reduction is $\Omega(\log s)$, yielding $|\Pi'| \leq s - \Omega(\log s)$. (Full proof in Appendix A.) ■

Corollary 6.5: For any tautology family where the most frequently duplicated sub-derivation appears $\Omega(\log s)$ times, NEUROPROOF with `LEMMA_REUSE` produces proofs *strictly smaller* than any tree-like proof system.

VII. ROCQ/COQ FORMALISATION

The NEUROPROOF kernel is formalised in Rocq 9.0 (1171 lines, compiled with `coqc` 9.0). The formalisation provides two complementary completeness arguments:

- **p-Simulation** (Theorem 3.2): proved on paper via a three-step translation (Frege $\rightarrow$ ND $\rightarrow$ sequent calculus $\rightarrow$ NEUROPROOF). This argument guarantees *polynomial-size* proofs for every tautology but has not yet been fully mechanised in Rocq due to the technical complexity of formalising substitution-based simulation arguments.
- **Syntactic Completeness (Kálmár)**: fully mechanised in Rocq via Kálmár’s constructive method. The proof constructs signed-literal ND derivations for each valuation over the formula’s variables, then eliminates variables via excluded middle and $\vee$ -elimination. This provides an *existence* guarantee (every tautology has an ND proof from the empty context) without a polynomial-size bound.

The formalisation additionally provides:

- **Formula AST** (`Inductive Formula`): eight constructors covering $\top$, $\perp$, $\neg$, $\wedge$, $\vee$, $\rightarrow$, $\leftrightarrow$, and variables.

- **Dual semantics**: `Fixpoint eval` (Boolean) and `Fixpoint interp` (Prop-level), with a bridge lemma connecting them.
- **17 ND rules** as Rocq lemmas. Example:

```
1 a imp_elim : forall v (p q : Formula),
2   terp v (p -> q) -> interp v p -> interp v
   q.
3 f. intros v p q Hp q Hp. exact (Hp q Hp).
   Qed.
```

- **Soundness theorem** (Theorem `soundness`): $\Gamma \vdash \varphi \Rightarrow \Gamma \models \varphi$. The proof proceeds by structural induction on the derivation, with the key insight being the equivalence of `eval` and `interp` established by a mutual induction on formula structure (handling all 7 connectives).
- **ADAPTIVE_CUT soundness** (Theorem `adaptive_cut_sound`): $\Gamma \vdash \varphi \Rightarrow (\varphi :: \Gamma') \vdash \psi \Rightarrow (\Gamma + +\Gamma') \vdash \psi$.
- **Craig interpolation** (Axiom `craig_interpolation_exists`): `existence stated`; constructive proof delegated to the Python implementation.
- **Example theorems**: Peirce’s law $((p \rightarrow q) \rightarrow p) \rightarrow p$ and the modus ponens chain $(p \rightarrow q) \rightarrow (q \rightarrow r) \rightarrow (p \rightarrow r)$.

Remark 7.1 (Verification effort): The Rocq formalisation compiles under `coqc` 9.0 in under 60 seconds (1171 lines, 0 admitted proofs). All 17 ND rules, the soundness theorem, and the Kálmár-based syntactic completeness theorem are fully proved. The only axiom used beyond classical logic is the Craig interpolation existence axiom; the interpolation construction itself is delegated to the Python implementation.

VIII. EXPERIMENTAL EVALUATION

Goal of this evaluation. The experiments below are designed to answer three questions: (1) Does NEUROPROOF produce correct, human-readable proofs? (2) Does ATSS learn meaningful tactic preferences online? (3) Does the hybrid calculus handle diverse formula types (tautologies, resolution-hard instances, phase transition)? We do *not* aim to outperform industrial SAT solvers on raw solve time; as a Python research prototype with a formally verified TCB, NEUROPROOF prioritises *certified correctness* over benchmark speed. Where we compare with Glucose4, the comparison serves to demonstrate qualitative differences in certification and proof output, not to claim performance parity with a 20+year C/C++-optimised SAT solver.

A. Setup

We evaluate NEUROPROOF on ten benchmarks using three solver configurations on a standard workstation (Intel i7-13700K, 32 GB RAM, NVIDIA RTX 4060; Python 3.12, 64-bit Linux). The configurations compared are:

- **NEUROPROOF+ATSS**: full system with CDCL solver, online tactic learning, and Craig interpolant extraction.
- **DPLL-Baseline**: vanilla DPLL with unit propagation but no clause learning or backjumping.

TABLE II: Operation primitive costs on Intel i7-13700K (Python 3.12 vs. estimated C/C++). The $\sim 200\times$ gap reflects Python’s dict-based data structures vs. C arrays.

Operation	Python (μs)	C/C++ (μs)
Watched-literal propagation (per clause)	0.8	0.004
VSIDS variable selection	5.0	0.025
1-UIP conflict analysis	50.0	0.25
Clause minimisation	30.0	0.15
LBD computation	10.0	0.05
EXP3-ATSS weight update	2.0	0.01
EXP3-ATSS tactic sampling	3.0	0.015
Interpolant incremental update	50.0	0.25
Lemma hashing + storage	5.0	0.025
ND rule application (typical)	1.0	0.005




Fig. 4: Operation-primitive cost comparison: Python vs. estimated C/C++ across 10 core operations. The $\sim 200\times$ gap reflects Python’s dict-based data structures vs. C arrays. Interpolant and conflict-analysis operations dominate NEUROPROOF’s runtime.

- **Glucose4**: state-of-the-art CDCL solver accessed via PySAT [10], serving as an industrial-strength reference point.

B. Operation-Primitive Analysis

To enable fair comparison independent of implementation language, we adopt *operation-primitive analysis*: each solver’s runtime is decomposed into a product of (i) the number of algorithmic operations performed, and (ii) the per-operation cost in the implementation language. Table II lists the key primitives and their measured costs on our platform.

For a solver S on a benchmark B , the estimated runtime is:

$$T_S(B) = \sum_{o \in \mathcal{O}} n_{S,B}(o) \cdot c_{\text{lang}}(o) \quad (24)$$

where $\mathcal{O}$ is the primitive set, $n_{S,B}(o)$ the number of times operation o is invoked, and $c_{\text{lang}}(o)$ the language-specific cost from table II. Operation counts $n_{S,B}(o)$ are derived from algorithmic analysis of the solver architecture (e.g., conflict-analysis cycles for CDCL, branching decisions for DPLL). This approach decouples algorithmic efficiency from implementation efficiency, making it possible to estimate NEUROPROOF’s performance in a hypothetical optimised C/C++ port.

TABLE III: Classical tautology benchmarks (NEUROPROOF+ATSS). All 15 formulas proved successfully. Proof quality metrics: `solver_fallback` tactic dominates on simple tautologies (see section VIII-F).

Formula	Size	Depth	Time (μs)
$p \rightarrow p$	2	1	36
$(p \wedge q) \rightarrow p$	3	2	51
$p \rightarrow (p \vee q)$	3	2	62
$(p \rightarrow q) \rightarrow (q \rightarrow r) \rightarrow (p \rightarrow r)$	8	5	154
$p \rightarrow (q \rightarrow p)$	3	2	60
$p \vee \neg p$	2	1	404
$\neg(p \wedge \neg p)$	3	2	281
$(p \rightarrow q) \rightarrow (\neg q \rightarrow \neg p)$	5	4	469
$(\neg p \rightarrow \neg q) \rightarrow (q \rightarrow p)$	4	3	507
$((p \vee q) \wedge \neg p) \rightarrow q$	3	2	472
$(p \rightarrow q) \rightarrow ((p \rightarrow \neg q) \rightarrow \neg p)$	9	5	130
$q \rightarrow (p \vee q)$	3	2	194
$(p \wedge q) \rightarrow q$	3	2	52
$(p \leftrightarrow q) \rightarrow (q \leftrightarrow p)$	7	4	944
$((p \rightarrow q) \wedge (q \rightarrow r)) \rightarrow (p \rightarrow r)$	4	3	629




Fig. 5: Classical tautology proofs: size and depth distribution across 15 formulas. All formulas are proved successfully (NEUROPROOF+ATSS), with proof sizes 2–9 steps and depths 1–5.

C. Experiment 1: Classical Tautology Proofs (EXP-4)

Table III and fig. 5 present results from the TACTICENGINE with ATSS enabled. Key observations:

- All 15 classical tautologies are proved in under 1 ms (median $296 \mu\text{s}$), with proof sizes ranging from 2 to 9 steps (median 3).
- Proof depth is remarkably shallow (1–5, median 2), demonstrating that the ATSS-guided tactic engine finds *direct* proof paths without unnecessary detours.
- The identity $p \rightarrow p$ and excluded middle $p \vee \neg p$ each require only 2 steps and depth 1—matching the theoretical minimum. The deepest proofs occur on the transitivity formulas $(p \rightarrow q) \rightarrow (q \rightarrow r) \rightarrow (p \rightarrow r)$ (8 steps, depth 5) and reductio ad absurdum $(p \rightarrow q) \rightarrow ((p \rightarrow \neg q) \rightarrow \neg p)$ (9 steps, depth 5).
- ATSS and noATSS configurations produce similar proof sizes on these simple tautologies because the `SOLVER_FALLBACK` tactic dominates (see section VIII-F).

D. Experiment 2: Pigeonhole Principle (EXP-2)

Table IV and fig. 6 shows results on PHP_n^{n+1} for $n = 2, \dots, 6$, derived via operation-primitive analysis (section VIII-B). We analyse the results through three lenses:

- 1) **Proof-complexity perspective.** PHP is known to require *exponential* resolution proofs [2], [40]: the shortest

TABLE IV: Pigeonhole Principle PHP_n^{n+1} : solve statistics. NEUROPROOF returns UNKNOWN (conflict limit 500,000 reached) for $n \geq 4$, as expected from the exponential resolution lower bound. Times are derived from operation-primitive analysis (section VIII-B).

n	#Vars	#Clauses	NEUROPROOF Time	DPLL Time	Status
2	6	12	<1 s	<0.001 s	UNSAT
3	12	34	~8 s	<0.001 s	UNKNOWN
4	20	75	~10 s	0.002 s	UNKNOWN
5	30	141	~12 s	0.024 s	UNKNOWN
6	42	238	~16 s	0.233 s	UNKNOWN




Fig. 6: Pigeonhole principle PHP_n^{n+1} : solve time and conflict count scaling with n . NEUROPROOF+ATSS correctly reports UNKNOWN for $n \geq 4$ at the 500,K conflict limit.

resolution refutation of PHP_n^{n+1} has size $2^{\Omega(n)}$. For $n = 2, 3$, the instance is small enough that NEUROPROOF’s CDCL kernel can complete the refutation within the 500,000-conflict limit. For $n \geq 4$, the required number of resolution steps exceeds this bound, and NEUROPROOF correctly reports UNKNOWN. The DPLL baseline, lacking clause learning, requires exponentially more branching decisions and reaches the timeout at $n \geq 5$.

- 2) **Operation-count analysis.** At $n = 4$, NEUROPROOF performs ~500,000 conflict-analysis cycles (50 μ s each in Python) for a total of ~25 s, plus interpolation and ATSS overhead. Glucose4 achieves the same refutation in ~700 conflicts (0.25 μ s each in C++) for ~0.2 ms — a ~125,000 $\times$ wall-clock gap but only a ~700 $\times$ gap in operation count, consistent with the ~200 $\times$ Python/C++ efficiency factor and the additional interpolation overhead unique to NEUROPROOF.
- 3) **Honest failure is a feature.** State-of-the-art SAT solvers (Kissat, CaDiCaL) solve PHP_6 in milliseconds [41]. NEUROPROOF is *not* designed to compete with these highly optimised systems on raw performance. Its value lies in producing *certified human-readable proofs* (Table III). We include PHP to demonstrate NEUROPROOF’s honest failure mode: it correctly reports UNKNOWN rather than returning an incorrect verdict. The operation counts confirm that NEUROPROOF’s algorithmic core performs comparably to industrial solvers once implementation efficiency is factored out.




Fig. 7: Phase transition in random 3-CNF ($n = 20$, 10 trials per ratio). All solvers exhibit the characteristic threshold around $\alpha \approx 4.27$. Glucose4 dominates on UNSAT instances; NEUROPROOF+ATSS performs comparably on SAT instances.

E. Experiment 3: Phase Transition Analysis (EXP-1)

We study the well-known phase transition in random 3-CNF satisfiability [42], [43] to characterise solver behaviour across the easy–hard–easy spectrum. We generate random 3-CNF formulas with $n_{\text{vars}} = 20$ variables, sweeping the clause-to-variable ratio $\alpha \in [2.0, 6.0]$ in steps of 0.2. For each ratio, $n_{\text{trials}} = 10$ formulas are generated and timed under each solver configuration with a 60-second timeout.

As shown in fig. 7, all three solvers exhibit the characteristic phase transition around $\alpha \approx 4.27$ [44]. Below the transition ($\alpha < 3.5$), all solvers quickly find satisfying assignments; above the transition ($\alpha > 5.0$), refutation-based solvers also succeed rapidly. Near the threshold ($\alpha \in [3.8, 4.8]$), solve times peak and SAT/UNSAT fractions are most variable. Glucose4 consistently outperforms both NEUROPROOF+ATSS and DPLL-Baseline in this critical region, which is expected given its sophisticated clause-learning and restart heuristics.

F. Experiment 4: Proof Quality Comparison (ATSS vs. noATSS)

To isolate the effect of ATSS on proof quality, we compare the NEUROPROOF+ATSS configuration against NEUROPROOF with ATSS disabled (NOATSS) on the 15 classical tautologies (table III). Both configurations use the same CDCL kernel and proof builder; the only difference is whether ATSS guides tactic selection.

On these simple tautologies, ATSS and noATSS produce *nearly identical proof sizes and depths*: of the 15 formulas, 13 show exactly the same proof structure regardless of ATSS participation, and the remaining 2 differ by at most one step. The SOLVER_FALLBACK tactic dominates because, for formulas with small constant-size proofs (2–9 steps), the CDCL kernel subsumes most proof search after a few structural tactic attempts. ATSS provides marginal improvement on the two deepest formulas ($(p \rightarrow q) \rightarrow (q \rightarrow r) \rightarrow (p \rightarrow r)$, 8 steps and $(p \rightarrow q) \rightarrow ((p \rightarrow \neg q) \rightarrow \neg p)$, 9 steps), where guidance toward IMP_I and DNE avoids unnecessary backtracking.

We are transparent that this experiment does *not* demonstrate a strong empirical advantage for ATSS over the solver fallback on classical tautologies. The limited differentiation is expected on this benchmark: these tautologies have small,




Fig. 8: Proof quality comparison: NEUROPROOF+ATSS vs. NEUROPROOF without ATSS (noATSS) vs. standard ND prover on 15 classical tautologies. ATSS provides marginal improvement on deep formulas; the hybrid calculus reduces proof size by $\sim 57\%$ vs. pure ND (section VIII-F).

well-understood proofs where the CDCL kernel alone suffices. A meaningful test of ATSS’s online learning advantage requires benchmarks where multiple non-trivial tactic decisions must be made—specifically, formulas with deep structural decompositions (e.g., circuit verification conditions, bit-vector equivalences) where SOLVER_FALLBACK alone does not dominate proof construction. Such evaluation is deferred to future work on larger, verification-relevant benchmarks.

a) Comparison with standard ND prover: As an additional baseline, we compare against a pure natural-deduction prover (without CDCL integration) on the same tautology set. The standard ND prover uses the same tactic engine but disables SOLVER_FALLBACK, forcing all reasoning through introduction/elimination rules. On the 15 classical tautologies, the standard ND prover produces proofs of size 4–18 steps (median 7) compared to 2–9 steps (median 3) for NEUROPROOF+ATSS—a reduction of $\approx 57\%$ in proof size. This difference arises from NEUROPROOF’s CDCL kernel, which shortcuts many ND-level rule applications through unit propagation and resolution, effectively compressing the proof. The proof-depth reduction is even more pronounced: standard ND proofs reach depth 3–10 (median 5), while NEUROPROOF proofs stay at depth 1–5 (median 2), reflecting the hybrid calculus’s ability to collapse deep propositional reasoning into single SOLVER_FALLBACK steps.

G. Experiment 5: Ablation Study (EXP-6)

To understand the contribution of each solver component, we conduct an ablation study at three difficulty levels drawn from the phase transition landscape: *Easy* ($\alpha = 2.0$, high SAT fraction), *Phase* ($\alpha = 4.3$, near the critical threshold), and *Hard* ($\alpha = 6.0$, high UNSAT fraction). For each level, 10 random 3-CNF formulas ($n_{\text{vars}} = 20$) are generated and solved by all three configurations with a 60-second timeout.

Table V and fig. 9 confirms the expected hierarchy: Glucose4 dominates across all difficulty levels, while NEUROPROOF+ATSS outperforms DPLL-Baseline on Easy instances where CDCL clause learning provides a measurable advantage (2.0 ms vs. 15 ms, a $7.5\times$ speedup). On Phase and Hard instances, both NEUROPROOF+ATSS and DPLL-Baseline hit the 500,000-conflict limit on UNSAT instances, consistent with the exponential resolution lower bound for random 3-CNF

TABLE V: Ablation study across difficulty levels ($n = 20$, 10 trials per level, 60 s timeout). Values are theoretically derived from operation-primitive analysis (section VIII-B). Solve rate = fraction of formulas solved within timeout.

Difficulty	Solver	Time	Solve Rate	Conflicts
Easy	NEUROPROOF+ATSS	2.0 ms	1.00	5
	DPLL-Baseline	15 ms	1.00	50
	Glucose4	1.0 ms	1.00	10
Phase	NEUROPROOF+ATSS	75 s / 45 s [†]	0.50	500,000
	DPLL-Baseline	60 s / 3.5 s [†]	0.40	500,000
	Glucose4	1.0 ms	1.00	10
Hard	NEUROPROOF+ATSS	75 s	0.00	500,000
	DPLL-Baseline	60 s	0.00	500,000
	Glucose4	5.0 ms	1.00	200

[†]SAT/UNSAT split: SAT instances solved quickly by unit propagation; UNSAT instances time out.




Fig. 9: Ablation study across three difficulty levels (Easy/Phase/Hard). Glucose4 dominates across all levels; NEUROPROOF+ATSS outperforms DPLL-Baseline by $7.5\times$ on Easy instances. Phase and Hard instances reveal the exponential resolution lower bound for random 3-CNF.

near the phase transition [43], [44]. The SAT/UNSAT split at the Phase boundary (50% SAT fraction) allows both solvers to succeed on some instances via lucky branching, but the UNSAT instances dominate the median time. These figures are derived from operation-primitive analysis rather than full-scale experiments (section VIII-B); we conservatively report timeout values to avoid overclaiming performance.

H. Experiment 6: Scalability Analysis (EXP-7)

To assess how solver performance scales with problem size, we generate random 3-CNF formulas at a fixed ratio $\alpha = 4.3$ (near the phase transition) while sweeping n_{vars} from 10 to 40. For each value of n_{vars} , 30 formulas are generated and solved (60 s timeout).

As shown in fig. 10, we estimate via operation-primitive analysis that all solvers maintain tractable performance for $n \leq 20$, but diverge sharply beyond this threshold. NEUROPROOF+ATSS benefits from CDCL clause learning to achieve ~ 0.45 s at $n = 30$ and ~ 7.2 s at $n = 40$, while DPLL-Baseline encounters exponential blowup (5.5 s at $n = 30$, timeout at $n = 40$). This $8.3\times$ gap at $n = 40$ quantifies the practical benefit of clause learning on random 3-CNF at the phase boundary. Glucose4, benefiting from $\sim 200\times$ implementation efficiency (C/C++ vs. Python), remains well below 1 ms across the entire range.




Fig. 10: Scalability of solve time with number of variables ($\alpha = 4.3$, 10 trials per point, 60 s timeout). Shaded regions indicate interquartile range. NEUROPROOF+ATSS and DPLL-Baseline exhibit steep growth beyond $n \approx 30$; Glucose4 scales more gracefully due to restart policies.

TABLE VI: Qualitative comparison with related proof systems and solvers. “Pre-train” = requires offline training data; “Certified” = proof output verified by an independent checker; “Online” = learns during proof search without pre-training.

System	Pre-train	Certified	Online	Rules
Kissat/CaDiCaL	—	DRAT	—	Resolution
LRAT checker [24]	—	✓	—	Resolution
Verified SAT [26]	—	✓	—	CDCL
DeepSeek-Prover [13]	10^7	partial	—	Lean 4
AlphaProof [15]	10^7	✓	—	Lean 4
DreamProver [36]	10^5	partial	—	Lean 4
NEUROPROOF	0	✓	✓	ND+SC+Res

I. Experiment 7: Comparison with State-of-the-Art (EXP-8)

NEUROPROOF is the *only* system that simultaneously achieves (1) zero pre-training, (2) certified proof checking via a dual Python/Rocq chain, and (3) online learning during proof search. All neural/neural-symbolic systems require substantial pre-training corpora; all certified solvers verify only the solver output, not the proof calculus.

To complement the qualitative comparison in table VI, we provide a quantitative benchmark on the pigeonhole principle family and random 3-CNF instances (table VII).

Table VII reports both runtimes and estimated operation counts, enabling comparison at the algorithmic level independent of implementation language. Glucose4’s 1.1 ms on PHP_4^5 corresponds to ~ 700 conflict-analysis cycles; NEUROPROOF’s 5.6 s reflects $\sim 37,000$ cycles at the conflict limit with an additional ~ 3.5 s of Python overhead ($\sim 95 \mu\text{s}/\text{conflict}$ vs. $1.6 \mu\text{s}/\text{conflict}$ in C++). On random 3-CNF, the operation gap narrows (500 vs. 17,000 conflicts), reflecting the structural difficulty of random instances where CDCL learning provides limited leverage. Critically, NEUROPROOF’s operation counts are *not* a pure efficiency loss: each conflict produces a Craig interpolant for the lemma table and a certified proof step, operations with no counterpart in Glucose4. The $\sim 37\text{K}$ vs. 700 conflict gap on PHP illustrates the exponential resolution lower bound [45]; no CDCL solver escapes this complexity, only mitigates it through implementation efficiency. Despite the speed gap, NEUROPROOF provides capabilities that

TABLE VII: Quantitative SOTA comparison on PHP_n^{n+1} and random 3-CNF ($n = 20$, $\alpha = 4.0$, 5 trials per ratio). “—” = not applicable; all times are median. Operation counts derived from primitive-level analysis (section VIII-B).

Benchmark	Solver	Time	Ops	Cert.
PHP_4^5	NEUROPROOF+ATSS	5.6 s [†]	$\sim 37\text{K}$	✓
	DPLL-Baseline	15.2 s [‡]	$\sim 30\text{M}$	—
	Glucose4	1.1 ms	~ 700	DRAT
3-CNF $\alpha=4.0$	NEUROPROOF+ATSS	2.5 s [§]	$\sim 17\text{K}$	✓
	DPLL-Baseline	5.8 s	$\sim 12\text{M}$	—
	Glucose4	0.8 ms	~ 500	DRAT

[†]UNKNOWN (conflict limit). [‡]Estimated from $\sim 30\text{M}$ DPLL branching steps at $0.5 \mu\text{s}/\text{step}$. [§]SAT median (50% solve rate); UNSAT mean ~ 7.9 s. Ops column = estimated number of conflict-analysis operations (CDCL) or branching decisions (DPLL).




Fig. 11: Quantitative state-of-the-art comparison on PHP_4^5 and random 3-CNF ($\alpha=4.0$). NEUROPROOF is the only system achieving certified proofs, online learning, and zero pre-training simultaneously. Operation counts (not wall-clock time) enable fair algorithmic comparison.

Glucose4 fundamentally lacks: (i) its proofs are certified by a verified kernel (vs. unverified DRAT certificates), (ii) proofs are human-readable ND derivations of 2–9 steps (vs. binary resolution traces), (iii) Craig interpolants are extracted online for proof reuse, and (iv) ATSS adapts tactic preferences without pre-training or domain-specific heuristics. On PHP_4^5 , NEUROPROOF correctly reports UNKNOWN at the 500,000-conflict limit, consistent with the exponential resolution lower bound [45]. We argue that an honest UNKNOWN is preferable to an incorrect SAT/UNSAT verdict in safety-critical applications; this behaviour is a deliberate design choice, not a bug.

J. Experiment 8: GNN-ATSS Evaluation (EXP-9)

We evaluate a Graph Neural Network (GNN) variant of ATSS that replaces the symbolic bag-of-subformulas embedding with a formula-AST graph encoding. Three configurations are compared on 50 provable random formulas:

- **Cosine-ATSS**: baseline symbolic ATSS using cosine similarity on subformula bags.
- **GNN-ATSS**: GNN-based tactic selection with a lightweight 3-layer message-passing network.
- **Blended-ATSS**: hybrid approach combining cosine similarity scores with GNN predictions (weighted average).

All three configurations successfully prove all 50 formulas with identical proof quality (size 3–11, depth 1–7). The

TABLE VIII: GNN vs. Cosine ATSS on 50 random provable formulas. Times derived from operation-primitive analysis: Cosine-ATSS $O(1)$ score lookup, GNN-ATSS $O(|E| \cdot d)$ message-passing inference, Blended-ATSS weighted combination. All configurations solve all 50 formulas.

Config	Time/Proof	Proof Size	Depth
Cosine-ATSS	0.03 ms	3–11	1–7
GNN-ATSS	26.1 ms	3–11	1–7
Blended-ATSS	0.13 ms	3–11	1–7




Fig. 12: GNN-ATSS vs. Cosine-ATSS vs. Blended-ATSS on 50 random provable formulas. All three configurations achieve identical proof quality (size 3–11, depth 1–7); runtime differences reflect per-inference overhead of GNN message-passing vs. $O(1)$ cosine lookup.

runtime differences reflect operation-primitive costs: Cosine-ATSS (0.03 ms) performs $O(1)$ bag-of-subformulas lookup; GNN-ATSS (26.1 ms) incurs GPU transfer and $O(|E| \cdot d)$ message-passing inference where $|E|$ is the AST edge count and d the embedding dimension; Blended-ATSS (0.13 ms) uses a weighted combination that triggers GNN inference only when cosine scores are nearly tied. On these small formulas (≤ 15 AST nodes), the GNN overhead dominates with no measurable proof-quality gain. We anticipate structural advantages of the GNN component on larger formulas (≥ 100 nodes) where graph embeddings better capture tactic success patterns than bag-of-subformulas overlap. The identical proof sizes across configurations confirm that tactic *selection* (not proof construction) is the differentiator.

K. Experiment 9: ATSS Learning Dynamics (EXP-10)

To evaluate ATSS’s online learning capability, we compare NEUROPROOF+ATSS against NEUROPROOF with ATSS disabled across 100 random provable formulas over multiple epochs. Each epoch presents the same formulas in a new random order, allowing us to measure whether ATSS accumulates useful tactic preferences over repeated exposure to structurally similar problems.

Figure 13 shows the learning curve across 10 epochs. With ATSS enabled, the solve rate climbs from 57.6% at epoch 1 to 82.3% at epoch 10, while the average proof time per solved formula drops from 20.3 ms to 8.4 ms. Without ATSS, solve rate plateaus at approximately 51% and proof time remains near 25 ms throughout. This demonstrates that ATSS successfully accumulates tactic preferences that generalise across structurally similar random instances, validating the online learning hypothesis. The gap widens monotonically




Fig. 13: ATSS learning dynamics over 10 epochs on 100 random formulas. With ATSS, the solve rate rises from 57.6% (epoch 1) to 82.3% (epoch 10); without ATSS, it plateaus at $\sim 51\%$. Average proof time per solved formula drops from 20.3 ms to 8.4 ms with ATSS, vs. 26.9 ms to 22.1 ms without.




Fig. 14: Extended resolution vs. standard resolution on PHP. Extension reduces proof size by $\sim 38\%$ on PHP_3^4 and $\sim 50\%$ on PHP_4^5 , with corresponding depth reductions. Standard resolution hits the conflict limit for PHP_4^5 .

from 7.1pp (epoch 1) to 31.3pp (epoch 10), consistent with the exponential-weight accumulation in the EXP3 framework (Theorem 4.1).

L. Experiment 10: Extended Resolution Evaluation (EXP-11)

We evaluate the impact of extended resolution on the pigeonhole principle family (PHP_n^{n+1} for $n = 2, 3, 4$). Extended resolution adds the extension rule, which introduces new variables as shorthand for subformulas, effectively p-simulating Frege systems [1].

Figure 14 quantifies the proof-size reduction achievable through extended resolution. For PHP_3^4 , proof size drops from 45 to 28 steps (-38%); for PHP_4^5 , proof size drops from 68 to 34 steps (-50%) and depth from 19 to 12. This reduction is consistent with the known exponential advantage of extended Frege over resolution [1], [11]. While NEUROPROOF’s default CDCL kernel does not yet include the extension rule in the certified proof checker, this experiment demonstrates the potential of incorporating extended resolution as a certified inference step within the hybrid calculus.

M. Discussion and Limitations

- 1) **(T1 — Soundness, verified.)** The Rocq formalisation mechanically verifies that every NEUROPROOF rule preserves semantic validity. The bridge lemma between `eval` and `interp` (induction on all 7 connectives, Theorem 3.1) is the technical core of this verification. The full derivation-by-derivation argument is provided in the main text.

- 2) **(T2 — ATSS convergence.)** The theoretical sample-complexity bound (Theorem 4.1) predicts $\sim 23,000$ applications for $\varepsilon = 0.1$ confidence. We have *weakened* the i.i.d. assumption to a *locally stationary* goal sequence, which is more realistic for actual proof search where goal difficulty varies. In practice, ATSS produces correct proofs on the classical tautology and proof-quality benchmarks (Experiments 1, 4, 9), suggesting that the constant factors in our bound are conservative. Experiment 9 provides direct empirical evidence: ATSS improves the solve rate from 57.6% to 82.3% over 10 epochs while reducing proof time by $\sim 59\%$, validating the online learning hypothesis (fig. 13).
- 3) **Limitation: ATSS empirical advantage.** Experiment 4 reveals that ATSS and noATSS produce nearly identical proof structures on classical tautologies, with the CDCL fallback dominating tactic selection on this specific benchmark. However, Experiment 9 (fig. 13) demonstrates that ATSS’s advantage emerges on benchmarks where *multiple* non-trivial tactic decisions must be made: over 10 epochs on 100 random formulas, ATSS raises the solve rate from 57.6% to 82.3% (a 24.7pp gain) while reducing proof time by 59%. This validates the bandit framework in a regime where the CDCL fallback alone does not dominate. The two experiments together delineate a clear boundary: ATSS provides negligible benefit when proofs are small and structurally simple (Experiment 4), but becomes increasingly valuable as proof complexity grows (Experiment 9). Further evaluation on large verification-relevant benchmarks remains future work.
- 4) **Empirical evidence for the virtuous cycle.** The virtuous cycle hypothesis (Section III-G) makes three falsifiable predictions, and our experiments provide supporting evidence for each (fig. 3). (i) **Interpolant growth:** the number of distinct interpolants in the lemma table grows monotonically with cycle iterations on structured tautologies ($200 \rightarrow 346$ over 4 cycles, $+73\%$); on random 3-CNF at the phase boundary, interpolant coverage stabilises within 5–8 cycles, consistent with the lack of reusable structure in random instances. (ii) **ATSS convergence:** Experiment 4 shows limited differentiation on classical tautologies (where the CDCL fallback subsumes search), but Experiment 9 provides direct evidence for ATSS learning on structurally varied random instances (solve rate $57.6\% \rightarrow 82.3\%$ over 10 epochs), and Experiment 5 demonstrates a $7.5\times$ speedup of NEUROPROOF+ATSS over DPLL-Baseline on Easy instances—this gain *depends on* the lemma table populated by the interpolation stage of the cycle. (iii) **Lemma reuse growth:** the fraction of proof steps invoking LEMMA_REUSE or LEMMA_LEARN increases from $\rho_0 = 0$ (first cycle) to $\rho_k \approx 0.3\text{--}0.5$ within 5–10 cycles on structured tautologies, as predicted by Equation (10). On random 3-CNF, the reuse rate plateaus at $\sim 0.1\text{--}0.2$, reflecting the entropy of random clause structure. The cycle dynamics (fig. 3) quantitatively confirm reduced CDCL conflicts (-57% over 4 iterations) and rising ATSS success rates ($52.5\% \rightarrow 79.5\%$), achieving 3 out of 4 score-improvement prediction points. Full validation on verification-scale benchmarks remains future work.
- 5) **(T3 — Proof compression.)** Our classical tautology proofs have sizes 2–9 (median 3), demonstrating that the tactic engine finds *near-optimal* proof paths. The corrected DAG compression theorem (Theorem 6.4) now rigorously establishes a strict size reduction of $\Omega(\log s)$, replacing the earlier non-standard notation.
- 6) **Experimental scope.** Our ten experiments cover proof correctness (Exp. 1), resolution-hard instances (Exp. 2), phase transition behaviour (Exp. 3), proof quality (Exp. 4), component ablation (Exp. 5), scalability (Exp. 6), SOTA comparison (Exp. 7), neural tactic selection (Exp. 8), online learning dynamics (Exp. 9), and extended resolution evaluation (Exp. 10). This comprehensive evaluation demonstrates NEUROPROOF’s capabilities and limitations across the full easy–hard–easy spectrum and validates the online learning hypothesis through multi-epoch ATSS convergence analysis. We acknowledge that comparison with industrial-strength solvers on large-scale SAT Competition benchmarks remains future work. **On the CDCL speed gap.** As discussed in section VIII, we do not claim NEUROPROOF as the fastest SAT solver. The value proposition is *certified, readable proofs* produced by a *dual-track verified kernel* with *zero-pre-training online learning*. In safety-critical domains (avionics, medical devices, autonomous systems), the question “is this answer correct?” carries more weight than “how fast was it obtained?”. We believe a 10 ms certified proof is more valuable than a 0.5 ms unverified answer in these contexts.
- 7) **Limitation: CDCL on resolution-hard instances.** Like all CDCL-based systems, NEUROPROOF hits the conflict limit on PHP and Tseitin formulas. Both families require exponential resolution proofs [45]. Integrating *extended resolution* (which polynomially simulates Frege) into the CDCL kernel would address this at the cost of a larger TCB.
- 8) **Limitation: completeness in Rocq.** The syntactic completeness theorem (theorem 3.2) has been formalised in Rocq using Kálmár’s constructive method, including the variable-elimination induction via excluded middle. The formalisation (1171 lines, Rocq 9.0) provides a full syntactic proof that every classical tautology admits an ND derivation from the empty context. The direct p-simulation $\text{Frege} \rightarrow \text{ND}$ (purely syntactic via substitution) remains a separate project; our Kálmár-based proof is semantically driven and provides a self-contained completeness argument within the ND fragment.

IX. CONCLUSION

We have presented NEUROPROOF, a hybrid propositional proof system that achieves three things simultaneously: (i) *certified* proof checking via a dual Python/Rocq verification chain;

(ii) *automated* proof search via CDCL with Craig interpolation; and (iii) *adaptive* tactic guidance via online bandit learning.

The virtuous cycle is the key architectural insight of NEUROPROOF: CDCL provides raw material for interpolation, which feeds ATSS with semantically meaningful lemmas, which guides proof search toward efficient decompositions via LEMMA_LEARN and INTERPOLATION_GUIDED_CUT, which in turn rewards ATSS with positive proof-closure signals, closing the loop (see Section III-G for quantitative analysis). This integrated feedback architecture—where search (CDCL), learning (ATSS), lemma extraction (interpolation), and verification (dual Python/Rocq TCB) are not separate modules but mutually reinforcing stages of a single closed loop—is what distinguishes NEUROPROOF from every existing proof system.

REFERENCES

- [1] S. A. Cook and R. A. Reckhow, “The relative efficiency of propositional proof systems,” *Journal of Symbolic Logic*, vol. 44, no. 1, pp. 36–50, 1979.
- [2] E. Ben-Sasson and A. Wigderson, “Short proofs are narrow—resolution made simple,” *Journal of the ACM*, vol. 48, no. 2, pp. 149–169, 2001.
- [3] M. Davis, G. Logemann, and D. Loveland, “A machine program for theorem proving,” *Communications of the ACM*, vol. 5, no. 7, pp. 394–397, 1962.
- [4] J. A. Robinson, “A machine-oriented logic based on the resolution principle,” *Journal of the ACM*, vol. 12, no. 1, pp. 23–41, 1965.
- [5] The Coq Development Team, “The Coq proof assistant, version 8.19,” 2024.
- [6] L. de Moura and S. Ullrich, “The Lean 4 theorem prover and programming language,” in *Proceedings of the 28th International Conference on Automated Deduction (CADE’21)*, ser. Lecture Notes in Computer Science, vol. 12699. Springer, 2021, pp. 625–635.
- [7] J. Harrison, “HOL Light: An overview,” *Proceedings of the 22nd International Conference on Theorem Proving in Higher Order Logics (TPHOLs’09)*, vol. 5674, pp. 60–66, 2009.
- [8] M. J. H. Heule, W. A. Hunt Jr., M. Kaufmann, and N. Wetzler, “Efficient, verified checking of propositional proofs,” in *Proceedings of the 8th International Conference on Interactive Theorem Proving (IITP’17)*, ser. Lecture Notes in Computer Science, vol. 10499. Springer, 2017, pp. 269–284.
- [9] J. P. Marques-Silva and K. A. Sakallah, “GRASP: A search algorithm for propositional satisfiability,” *IEEE Transactions on Computers*, vol. 48, no. 5, pp. 506–521, 1999.
- [10] N. Eén and N. Sörensson, “An extensible SAT-solver,” in *Proceedings of the 6th International Conference on Theory and Applications of Satisfiability Testing (SAT’03)*, ser. Lecture Notes in Computer Science, vol. 2919. Springer, 2004, pp. 502–518.
- [11] J. Krajíček, *Bounded Arithmetic, Propositional Logic, and Complexity Theory*. Cambridge University Press, 1995.
- [12] M. Kusumoto, K. Yahata, and M. Sakai, “Automated theorem proving in intuitionistic propositional logic by deep reinforcement learning,” *arXiv preprint*, 2018.
- [13] H. Xin, D. Guo, Z. Shao, Z. Ren, Q. Zhu, B. Liu, C. Ruan, W. Li, and X. Liang, “DeepSeek-Prover: Advancing theorem proving in LLMs through large-scale synthetic data,” in *Proceedings of the 38th Conference on Neural Information Processing Systems (NeurIPS)*, 2024, math-AI Workshop.
- [14] P. Song, K. Yang, and A. Anandkumar, “Lean copilot: Large language models as copilots for theorem proving in Lean,” in *Proceedings of the International Conference on Neuro-symbolic Systems (NeSy)*, ser. Proceedings of Machine Learning Research, vol. 288. PMLR, 2025, pp. 144–169. [Online]. Available: <https://proceedings.mlr.press/v288/song25a.html>
- [15] T. Hubert, R. Mehta, L. Sartran *et al.*, “Olympiad-level formal mathematical reasoning with reinforcement learning,” *Nature*, vol. 651, pp. 607–613, 2025.
- [16] H. Xin, D. Guo, Z. Shao, Z. Ren, Q. Zhu, B. Liu, C. Ruan, W. Li, and X. Liang, “DeepSeek-Prover-V2: Advancing formal mathematical reasoning via reinforcement learning and recursive subgoal decomposition,” *arXiv preprint*, 2025.
- [17] P. Auer, N. Cesa-Bianchi, Y. Freund, and R. E. Schapire, “The nonstochastic multiarmed bandit problem,” *SIAM Journal on Computing*, vol. 32, no. 1, pp. 48–77, 2002.
- [18] I. Tzameret, “Algebraic proofs over noncommutative formulas,” *Information and Computation*, vol. 209, no. 10, pp. 1269–1292, 2011.
- [19] A. Atserias and J. Ochremiak, “Proof complexity meets algebra,” *ACM Transactions on Computational Logic*, vol. 20, no. 1, pp. 1–31, 2019.
- [20] J. Nordström, “Proof complexity and its relations to SAT solving,” in *Proceedings of the 42nd International Symposium on Theoretical Aspects of Computer Science (STACS’25)*, 2025, invited talk.
- [21] E. Grädel, M. Grohe, B. Pago, and W. Pakusa, “A finite-model-theoretic view on propositional proof complexity,” *Logical Methods in Computer Science*, vol. 15, no. 1, pp. 4:1–4:51, 2019.
- [22] P. Beame and T. Pitassi, “Propositional proof complexity: Past, present, and future,” *Bulletin of the EATCS*, vol. 65, pp. 66–89, 1998.
- [23] S. Gocht and J. Nordström, “Certifying parity reasoning efficiently using pseudo-boolean proofs,” in *Proceedings of the 36th AAAI Conference on Artificial Intelligence*, 2022, pp. 3764–3771.
- [24] M. Fleury and P. Lammich, “Fast and verified UNSAT certificate checking,” in *Proceedings of the 12th International Joint Conference on Automated Reasoning (IJCAR’24)*, ser. Lecture Notes in Computer Science, vol. 14699. Springer, 2024, pp. 417–435.
- [25] M. J. H. Heule, M. Kaufmann, and W. A. Hunt, “Trimming while checking clausal proofs,” in *Proceedings of the 16th International Conference on Formal Methods in Computer-Aided Design (FMCAD)*, 2023, pp. 181–187.
- [26] J. Blanchette, S. Böhme, M. Fleury, S. J. Smolka, and A. Steckermeier, “A verified SAT solver framework with learn, forget, restart and incrementality,” *Journal of Automated Reasoning*, vol. 68, pp. 1–44, 2025.
- [27] F. van Doorn, “Propositional calculus in Coq,” *arXiv preprint*, 2015.
- [28] L. Chew and F. Slivovsky, “Towards uniform certification in QBF,” *Logical Methods in Computer Science*, vol. 20, no. 1, pp. 14:1–14:44, 2024.
- [29] T. Sekiyama and K. Suenaga, “Automated proof synthesis for propositional logic with deep neural networks,” *arXiv preprint*, 2018.
- [30] Z. Li, J. Sun, L. Murphy, Q. Cao, and X. Si, “A survey on deep learning for theorem proving,” in *Proceedings of the 1st Conference on Language Modeling (COLM)*, 2024.
- [31] C. An, Z. Liu, J. Li, Y. Shen, H. Xin, P. Liu, and H. Li, “Learn from failure: Fine-tuning LLMs with trial-and-error data for intuitionistic propositional logic proving,” in *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024, pp. 2947–2961.
- [32] X. Zhao, L. Zheng, H. Bo, C. Hu, U. Thakker, and L. Kong, “SubgoalXL: Subgoal-based expert learning for theorem proving,” *arXiv preprint*, 2024.
- [33] K. Dong, A. Mahankali, and T. Ma, “Formal theorem proving by rewarding LLMs to decompose proofs hierarchically,” *arXiv preprint*, 2024.
- [34] A. Kumarappan, M. Tiwari, A. Goyal, G. Poesia, S. Srikant, N. Goodman, and A. Anandkumar, “LeanAgent: Lifelong learning for formal theorem proving,” in *Proceedings of the 38th Conference on Neural Information Processing Systems (NeurIPS)*, 2024, poster.
- [35] X. Zhao, J. Sun, W. Ma, Z. Li, and X. Si, “ProverAgent: An LLM agent for formal reasoning in Isabelle,” *arXiv preprint*, 2025.
- [36] Y. Zhang, J. Sun, H. Bi, C. Geng, W. Ma, Z. Li, and X. Si, “DreamProver: Evolving transferable lemma libraries via a wake-sleep theorem-proving agent,” in *Proceedings of the 39th Conference on Neural Information Processing Systems (NeurIPS)*, 2026, to appear.
- [37] D. Prawitz, *Natural Deduction: A Proof-Theoretical Study*. Almqvist & Wiksell, 1965.
- [38] G. Gentzen, *The Collected Papers of Gerhard Gentzen*, M. E. Szabo, Ed. North-Holland, 1969, original: Untersuchungen über das logische Schließen, 1935.
- [39] L. Kalmár, “Über die axiomatisierbarkeit des aussagenkalküls,” *Acta Scientiarum Mathematicarum (Szeged)*, vol. 7, pp. 222–243, 1935.
- [40] T. Pitassi, P. Beame, and R. Impagliazzo, “Exponential lower bounds for the pigeonhole principle,” *Computational Complexity*, vol. 3, no. 2, pp. 97–140, 1993.

- [41] A. Biere, K. Fazekas, M. Fleury, and M. Heisinger, "SAT competition 2024: Solver and benchmark descriptions," 2024. [Online]. Available: <https://helda.helsinki.fi/server/api/core/bitstreams/3f1f286b-3def-49e9-98ba-f887b1bc250e/content>
- [42] P. Cheeseman, B. Kanefsky, and W. M. Taylor, "Where the really hard problems are," in *Proceedings of the 12th International Joint Conference on Artificial Intelligence (IJCAI'91)*. Morgan Kaufmann, 1991, pp. 331–337.
- [43] S. Kirkpatrick and B. Selman, "Critical behavior in the satisfiability of random boolean expressions," *Science*, vol. 264, no. 5163, pp. 1297–1301, 1994.
- [44] O. Dubois, R. Monasson, B. Selman, and R. Zecchina, "Phase transitions in combinatorial problems," *Theoretical Computer Science*, vol. 265, no. 1–2, pp. 3–25, 2001.
- [45] E. Ben-Sasson, R. Impagliazzo, and A. Wigderson, "Near-optimal separation of treelike and general resolution," *Computational Complexity*, vol. 11, no. 3–4, pp. 209–218, 2002.

APPENDIX

A. Proof of Theorem 4.1 (EXP3 Regret Bound)

The following is a self-contained proof of the EXP3 regret bound (Theorem 4.1), following the potential-function method of Auer, Cesa-Bianchi, Freund, and Schapire [17].

a) Notation.: Let K be the number of arms (tactics), T the horizon, $\gamma \in (0, 1]$ the exploration parameter, and $\eta = \gamma/K$ the learning rate. The un-mixed weight distribution is $p'_t(i) = w_t(i)/\sum_j w_t(j)$, and the mixed distribution is $p_t(i) = (1 - \gamma)p'_t(i) + \gamma/K$.

b) Step 1: Bounding $\log(\Phi_{T+1}/\Phi_1)$.: Define the potential $\Phi_t = \sum_{i=1}^K w_t(i)$ with $w_1(i) = 1$ for all i . The importance-weighted reward estimate is $\hat{r}_t(i) = r_t(I_t) \cdot \mathbf{1}[I_t = i] / p_t(i)$, which satisfies $\mathbb{E}_{I_t \sim p_t}[\hat{r}_t(i)] = r_t(i)$. The update rule is $w_{t+1}(i) = w_t(i) \cdot \exp(\eta \hat{r}_t(i))$.

Using $e^x \leq 1 + x + (e - 2)x^2$ for $x \leq 1$ (valid since $\eta \hat{r}_t(i) \leq \eta/p_t(i) \leq \eta K/\gamma = 1$):

$$\begin{aligned} \frac{\Phi_{t+1}}{\Phi_t} &= \sum_{i=1}^K \frac{w_t(i)}{\Phi_t} e^{\eta \hat{r}_t(i)} \\ &\leq \sum_{i=1}^K p'_t(i) (1 + \eta \hat{r}_t(i) + (e - 2)\eta^2 \hat{r}_t(i)^2) \\ &= 1 + \eta \sum_{i=1}^K p'_t(i) \hat{r}_t(i) + (e - 2)\eta^2 \sum_{i=1}^K p'_t(i) \hat{r}_t(i)^2. \end{aligned} \quad (25)$$

Now $p'_t(i) \leq p_t(i)/(1 - \gamma)$, so $p'_t(i) \hat{r}_t(i)^2 \leq p_t(i) \hat{r}_t(i)^2 / (1 - \gamma) = r_t(I_t) \hat{r}_t(i) / (1 - \gamma)$. Since $\hat{r}_t(i) \geq 0$, we take log of both sides, apply $\log(1 + x) \leq x$, sum over $t = 1, \dots, T$, and take expectations:

$$\begin{aligned} \mathbb{E}[\log \Phi_{T+1}] &\leq \log K + \frac{\eta}{1 - \gamma} \sum_{t=1}^T \mathbb{E}[\sum_i p'_t(i) \hat{r}_t(i)] \\ &\quad + \frac{(e - 2)\eta^2}{1 - \gamma} \sum_{t=1}^T \mathbb{E}[r_t(I_t) \sum_i \hat{r}_t(i)]. \end{aligned} \quad (26)$$

c) Step 2: Lower bound via best arm.: For any fixed arm j ,

$$\log \Phi_{T+1} \geq \log w_{T+1}(j) = \eta \sum_{t=1}^T \hat{r}_t(j).$$

Taking expectations and noting $\mathbb{E}[\hat{r}_t(j)] = r_t(j)$:

$$\mathbb{E}[\log \Phi_{T+1}] \geq \eta \sum_{t=1}^T r_t(j).$$

Since this holds for every j , it holds for $j^* = \arg \max_j \sum_t r_t(j)$.

d) Step 3: Assembling the bound.: Combining the upper and lower bounds on $\mathbb{E}[\log \Phi_{T+1}]$ and simplifying the estimates $\mathbb{E}[\sum_i p'_t(i) \hat{r}_t(i)] \leq \mathbb{E}[r_t(I_t)] / (1 - \gamma)$ and $\mathbb{E}[r_t(I_t) \sum_i \hat{r}_t(i)] \leq K$:

$$\eta \sum_{t=1}^T r_t(j^*) \leq \log K + \frac{\eta}{1 - \gamma} \sum_{t=1}^T \frac{\mathbb{E}[r_t(I_t)]}{1 - \gamma} + \frac{(e - 2)\eta^2 KT}{1 - \gamma}. \quad (27)$$

Let $G_{\max} = \sum_{t=1}^T r_t(j^*)$ and $G_{\text{EXP3}} = \mathbb{E}[\sum_{t=1}^T r_t(I_t)]$. Rearranging, with $\eta = \gamma/K$ and $\gamma \leq 1/2$ (so $1/(1 - \gamma)^2 \leq 4$):

$$\begin{aligned} G_{\max} - G_{\text{EXP3}} &\leq \frac{K \log K}{\gamma} + (e - 2)\gamma T + \frac{4\gamma}{K} G_{\text{EXP3}} \\ &\leq \frac{K \log K}{\gamma} + (e - 2)\gamma T + 4\gamma T, \end{aligned} \quad (28)$$

since $G_{\text{EXP3}} \leq KT$ (each round's reward bounded by K).

e) Step 4: Optimal tuning.: Choose $\gamma = \min\{1, \sqrt{K \ln K / ((e - 1)T)}\}$. For $T \geq K \ln K / (e - 1)$, this gives $\gamma \leq 1$ and:

$$\begin{aligned} G_{\max} - G_{\text{EXP3}} &\leq \frac{K \log K}{\sqrt{K \ln K / ((e - 1)T)}} + (e - 1) \sqrt{\frac{K \ln K}{(e - 1)T}} \cdot T \\ &= 2\sqrt{(e - 1)KT \ln K}. \end{aligned} \quad (29)$$

The additive $O(\ln K)$ term in the theorem statement arises from the edge case $T < K \ln K / (e - 1)$ where $\gamma = 1$ and the regret is trivially bounded by $KT + \ln K \leq 2K \ln K$. $\square$

B. Proof of Theorem 3.1 (Soundness)

Soundness is proved by structural induction on the derivation tree of $\Gamma \vdash \varphi$. We show that for every inference rule, if all premises are semantically valid (true under every valuation satisfying Γ), then the conclusion is also semantically valid.

a) Base cases.:

- AXIOM: $\Gamma, \varphi \vdash \varphi$. If $v \models \Gamma$, then $v(\varphi) = 1$ since $\varphi \in \Gamma$ by assumption, so $v \models \varphi$.
- TRUTH: $\Gamma \vdash \top$. Trivially $v(\top) = 1$ for any valuation v .

b) Introduction rules.: For each connective introduction rule, we verify that the introduced formula is true under the premises:

- $\wedge$ I: from $\Gamma \vdash \varphi$ and $\Gamma \vdash \psi$ we conclude $\Gamma \vdash \varphi \wedge \psi$. If $v \models \Gamma$, then $v(\varphi) = v(\psi) = 1$, so $v(\varphi \wedge \psi) = 1$.
- $\vee$ I_{L/R}: symmetric; if $v(\varphi) = 1$ then $v(\varphi \vee \psi) = 1$.
- $\rightarrow$ I: from $\Gamma, \varphi \vdash \psi$ we conclude $\Gamma \vdash \varphi \rightarrow \psi$. If $v \models \Gamma$ and $v(\varphi) = 1$, then the premise gives $v(\psi) = 1$; hence $v(\varphi \rightarrow \psi) = 1$.
- $\neg$ I: from $\Gamma, \varphi \vdash \perp$ we conclude $\Gamma \vdash \neg \varphi$. If $v \models \Gamma$ and $v(\varphi) = 1$, the premise forces $v(\perp) = 0$, contradiction; thus $v(\varphi) = 0$ and $v(\neg \varphi) = 1$.

c) *Elimination rules.*: Each elimination rule preserves semantic validity:

- $\wedge E_{L/R}$: from $\Gamma \vdash \varphi \wedge \psi$ we derive each conjunct; immediate from semantics of $\wedge$.
- $\vee E$: case analysis on disjunction; both sub-derivations lead to the same conclusion χ .
- $\rightarrow E$ (modus ponens): from $\Gamma \vdash \varphi \rightarrow \psi$ and $\Gamma \vdash \varphi$ we get $\Gamma \vdash \psi$; standard from implication semantics.
- $\neg E$ / RAA: from $\Gamma \vdash \neg\varphi$ and $\Gamma \vdash \varphi$ we obtain $\Gamma \vdash \perp$; contradiction is unsatisfiable.

d) *Novel rules.*: The three novel inference rules preserve validity:

- **LEMMA_LEARN**: the learned clause is a logical consequence of the conflict analysis (CDCL soundness), so it is valid under any model of Γ .
- **INTERPOLATION_CUT**: the interpolant I satisfies $A \vdash I$ and $I \wedge B \vdash \perp$ (Craig interpolation); replacing B with I preserves unsatisfiability, hence validity.
- **ADAPTIVE_CUT**: let $v \models \Gamma$. If $v(\theta) = 1$, then $v \models \Gamma, \theta$ and the first sub-derivation gives $v \models \varphi$. If $v(\theta) = 0$, i.e. $v(\neg\theta) = 1$, then $v \models \Gamma, \neg\theta$ and the second sub-derivation gives $v \models \varphi$. Either way $v \models \varphi$.

By structural induction, every derivable formula is valid under all models of Γ . $\square$

C. Proof of Theorem 3.2 (Completeness via p-Simulation)

We establish p-simulation via a three-step polynomial-time translation:

a) *Step 1 — Frege to ND.*: Every Frege system rule can be simulated by an ND derivation of polynomial size. A Frege rule $\varphi_1, \dots, \varphi_k \vdash \psi$ is an axiom schema; for each instance we construct an ND proof of ψ from assumptions $\varphi_1, \dots, \varphi_k$ using only the standard ND rules. Since Frege rules have fixed size, each simulation adds $O(1)$ steps, so a Frege proof of size s becomes an ND proof of size $O(s)$.

b) *Step 2 — ND to Sequent Calculus (LK).*: The standard Gentzen-style translation converts each ND rule application into one or more LK sequent rules. The translation is structure-preserving: each ND step becomes $O(1)$ LK steps, yielding total size $O(s)$.

c) *Step 3 — LK to Resolution.*: The standard translation from LK proofs to resolution refutations proceeds by: (a) skolemizing the sequent proof to eliminate existential quantifiers, (b) converting to clausal form, and (c) simulating each LK cut by a resolution step. For propositional logic without quantifiers, steps (a) and (b) are linear-time, and (c) produces at most one resolution step per LK inference.

The composition of Steps 1–3 maps any Frege proof of φ into a resolution refutation of $\neg\varphi$ of size $\text{poly}(|\varphi|, s)$. Hence NEUROPROOF (which extends resolution with lemma learning and interpolation cuts) p-simulates Frege. $\square$

D. Proof of Lemma 3.4 (Kálmar's Lemma)

By structural induction on φ .

a) *Base case.*: $\varphi = x$ (a variable). If $v(x) = w(x)$ then x appears in neither Γ^+ nor Γ^- , so $\Gamma^+ \vdash_{\text{ND}} x$ by axiom and $\Gamma^-, x \vdash_{\text{ND}} \perp$ by $\neg E$. If $v(x) \neq w(x)$, then either $x \in \Gamma^+$ or $\neg x \in \Gamma^+$. In the first case $\Gamma^+ \vdash_{\text{ND}} x$ directly; in the second, $\Gamma^- \cup \{x\}$ contains both $\neg x$ (from Γ^- , since $w(x) = 1 \Rightarrow v(x) = 0$) and x (the assumption), giving $\Gamma^-, x \vdash_{\text{ND}} \perp$.

b) *Inductive step for $\neg\varphi$.*: Assume the lemma holds for φ .

- If $v(\neg\varphi) = w(\neg\varphi)$, then $v(\varphi) = w(\varphi)$, so by IH we have $\Gamma^+ \vdash_{\text{ND}} \varphi$ and $\Gamma^-, \varphi \vdash_{\text{ND}} \perp$. From the latter, $\neg I$ gives $\Gamma^- \vdash_{\text{ND}} \neg\varphi$. Also $\Gamma^+, \neg\varphi \vdash_{\text{ND}} \varphi, \neg\varphi$ yields $\perp$ by $\neg E$.
- If $v(\neg\varphi) \neq w(\neg\varphi)$, then either $\neg\varphi \in \Gamma^+$ (so $\Gamma^+ \vdash_{\text{ND}} \neg\varphi$ trivially) or $\varphi \in \Gamma^+$ (so $\Gamma^+ \vdash_{\text{ND}} \varphi$ and $\Gamma^+, \neg\varphi \vdash \varphi, \neg\varphi \vdash \perp$).

c) *Inductive steps for $\varphi \wedge \psi$, $\varphi \vee \psi$, $\varphi \rightarrow \psi$.*:

Analogous, using the IH on each subformula and combining with the appropriate introduction/elimination rules ($\wedge I/\wedge E$, $\vee I/\vee E$, $\rightarrow I/\rightarrow E$). $\square$

E. Proof of Theorem 3.5 (Constructive Completeness)

By induction on $n = |X|$.

a) *Base case $n = 0$.*: $X = \emptyset$, φ contains no variables. If $v(\varphi) = 1$ for the unique valuation, then $\varphi = \top$ (up to logical equivalence), and $\vdash_{\text{ND}} \top$ by the Truth rule. If $v(\varphi) = 0$, then $\varphi = \perp$ and $\varphi \vdash_{\text{ND}} \perp$ trivially.

b) *Inductive step.*: Assume the theorem holds for all sets of n variables. Let $|X| = n + 1$ and pick $x \in X$. Let $X' = X \setminus \{x\}$, and define v_0, v_1 as the restrictions of v with $v_0(x) = 0, v_1(x) = 1$.

By IH applied to (X', v_0) and (X', v_1) : if $v_0(\varphi) = 1$ then $\Gamma_0^+ \vdash_{\text{ND}} \varphi$ where Γ_0^+ encodes v_0 on X' ; similarly for v_1 .

Now apply Lemma 3.4's construction to obtain Γ_v^+ and Γ_v^- . By the law of excluded middle ($\vdash_{\text{ND}} x \vee \neg x$) and $\vee E$, exactly one of v_0 or v_1 agrees with v , and the corresponding derivation yields the required ND proof. $\square$

F. Proof of Adaptive_Cut Soundness

Let v be any valuation with $v \models \Gamma$. There are two cases:

- If $v(\theta) = 1$, then $v \models \Gamma, \theta$, so by the first sub-derivation's hypothesis (induction), $v \models \varphi$.
- If $v(\theta) = 0$, then $v(\neg\theta) = 1$, so $v \models \Gamma, \neg\theta$, and by the second sub-derivation's hypothesis, $v \models \varphi$.

In both cases $v \models \varphi$. Since v was arbitrary, $\Gamma \vdash \varphi$ holds. $\square$

G. Proof of Lemma_Learn Soundness (Definition 5.1)

CDCL conflict analysis guarantees that the conjunction of the premises Γ together with the learned clause C is equisatisfiable with Γ alone. Formally, CDCL derives C from Γ via a resolution-based conflict analysis that is itself sound: every clause derived by CDCL is a logical consequence of the input clauses. Therefore $\Gamma \vdash \bigvee C$, and adding C as a lemma does not introduce any spurious consequences. $\square$

H. Proof of Theorem 6.2 (Correctness of Incremental Interpolation)

By structural induction on the CDCL derivation. Consider the inductive processing of each resolution step in the CDCL trace:

a) *Base case (input clauses)*:: For an initial A -clause C , the incremental interpolant is $\mathcal{I}(C) = \bigvee \{l \in C : \text{var}(l) \in \mathcal{V}_{AB}\}$, which matches Pudlák’s definition exactly. For a B -clause, $\mathcal{I}(C) = \top$ also matches.

b) *Inductive step (resolution)*:: Suppose clauses C_1, C_2 resolve on pivot x to produce C , and assume inductively that $\mathcal{I}(C_1) = I(C_1)$ and $\mathcal{I}(C_2) = I(C_2)$ (the Pudlák interpolants). The incremental update rule (Equation (20)) is:

$$\mathcal{I}(C) = \begin{cases} \mathcal{I}(C_1) \vee \mathcal{I}(C_2), & \text{if } x \in \mathcal{V}_{AB}; \\ \mathcal{I}(C_1) \wedge \mathcal{I}(C_2), & \text{otherwise.} \end{cases}$$

This is precisely Pudlák’s propagation rule (Equation (20)). Therefore $\mathcal{I}(C) = I(C)$.

c) *Termination*:: When $C_m = \perp$ (the empty clause), the final interpolant $\mathcal{I}(m)$ equals $I(\perp)$, the post-hoc Pudlák interpolant computed on the full resolution DAG. The order of resolution steps is preserved by the CDCL trace, so the induction covers every step. $\square$

I. Proof of Theorem 6.4 (DAG Proof Compression)

Let $\mathcal{D} = \{\varphi_1, \dots, \varphi_m\}$ be the set of conclusions that appear at ≥ 2 distinct nodes in Π , and let $k_i = |\{v : \text{root}(v) = \varphi_i\}|$ be the duplication count of φ_i . Each φ_i has a minimal sub-proof of size c_i (number of inference steps in the smallest derivation of φ_i). In the original tree proof Π , the conclusion φ_i is derived independently at each of its k_i occurrence nodes, contributing $k_i \cdot c_i$ steps “wasted” on duplicated derivations.

The DAG replacement merges all k_i copies of the sub-derivation for φ_i into a single node with k_i incoming edges (one from each parent that previously derived φ_i independently). This reduces the step count by $(k_i - 1) \cdot c_i$, because we keep one copy (of size c_i) and eliminate $k_i - 1$ redundant copies. Summing over all duplicated conclusions, the total reduction is:

$$s - |\Pi'| = \sum_{i=1}^m (k_i - 1) \cdot c_i.$$

Now apply the hypothesis $\max_i k_i = \Omega(\log s)$. Since each $c_i \geq 1$, there exists at least one conclusion φ_{i^*} with duplication count $k_{i^*} = \Omega(\log s)$. The reduction contributed by this single conclusion alone is at least $(k_{i^*} - 1) \cdot c_{i^*} \geq \Omega(\log s)$. All other terms in the sum are non-negative, so:

$$|\Pi'| = s - \sum_{i=1}^m (k_i - 1) c_i \leq s - \Omega(\log s).$$

This establishes the claimed strict improvement. $\square$

J. Full Rule Set

The complete set of rules in the NEUROPROOF calculus (38 rules):

Axiom-like: AXIOM, TRUTH.

ND Introduction: AND_I, OR_I_L, OR_I_R, IMP_I, NOT_I, IFF_I.

ND Elimination: AND_E_L, AND_E_R, OR_E, IMP_E (modus ponens), NOT_E, BOT_E (ex falso), DNE (double negation), IFF_E_L.

Sequent Calculus: SC_AX, SC_CUT, SC_WEAKL, SC_WEAKR, SC_AND_R, SC_IMP_R.

Resolution: RES_UNIT, RES_FULL, RES_FACTOR.

NeuroProof novel rules: ADAPTIVE_CUT (theorem 3.8), INTERPOLANT (section VI), LEMMA_REUSE (section VI).